

Amazon Fine Food Reviews Analysis

Data Source: <https://www.kaggle.com/snap/amazon-fine-food-reviews>

EDA: <https://nycdatascience.com/blog/student-works/amazon-fine-foods-visualization/>

The Amazon Fine Food Reviews dataset consists of reviews of fine foods from Amazon.

Number of reviews: 568,454

Number of users: 256,059

Number of products: 74,258

Timespan: Oct 1999 - Oct 2012

Number of Attributes/Columns in data: 10

Attribute Information:

1. Id
2. ProductId - unique identifier for the product
3. UserId - unique identifier for the user
4. ProfileName
5. HelpfulnessNumerator - number of users who found the review helpful
6. HelpfulnessDenominator - number of users who indicated whether they found the review helpful or not
7. Score - rating between 1 and 5
8. Time - timestamp for the review
9. Summary - brief summary of the review
10. Text - text of the review

Objective:

Given a review, determine whether the review is positive (rating of 4 or 5) or negative (rating of 1 or 2).

[Q] How to determine if a review is positive or negative?

[Ans] We could use Score/Rating. A rating of 4 or 5 can be considered as a positive review. A rating of 1 or 2 can be considered as negative one. A review of rating 3 is considered neutral and such reviews are ignored from our analysis. This is an approximate and proxy way of determining the polarity (positivity/negativity) of a review.

[1]. Reading Data

[1.1] Loading the data

The dataset is available in two forms

1. .csv file
2. SQLite Database

In order to load the data, We have used the SQLITE dataset as it is easier to query the data and visualise the data efficiently.

Here as we only want to get the global sentiment of the recommendations (positive or negative), we will purposefully ignore all Scores equal to 3. If the score is above 3, then the recommendation will be set to "positive". Otherwise, it will be set to "negative".

```
In [0]: %matplotlib inline
import warnings
warnings.filterwarnings("ignore")

import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
```

```

import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer

import re
# Tutorial about Python regular expressions: https://pymotw.com/2/re/
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle

from tqdm import tqdm
import os

```

```

In [3]: from google.colab import drive
drive.mount('/content/drive')

```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```

In [0]: !cp "/content/drive/My Drive/final.sqlite" "final.sqlite"

```

```

In [5]: import os
if os.path.isfile('final.sqlite'):
    conn = sqlite3.connect('final.sqlite')
    final = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score !=

```

```

3 """ , conn)
    conn.close()
else:
    print("Please the above cell")

print("Preprocessed Amzon fine food data columns shape : ",final.shape
)
print("fPreprocessed Amzon fine food data columns      :",final.column
s.values)

```

```

Preprocessed Amzon fine food data columns shape : (364171, 12)
fPreprocessed Amzon fine food data columns      : ['index' 'Id' 'Produ
ctId' 'UserId' 'ProfileName' 'HelpfulnessNumerator'
'HelpfulnessDenominator' 'Score' 'Time' 'Summary' 'Text' 'CleanedTex
t']

```

```

In [0]: # using SQLite Table to read data.
con = sqlite3.connect('database.sqlite')

# filtering only positive and negative reviews i.e.
# not taking into consideration those reviews with Score=3
# SELECT * FROM Reviews WHERE Score != 3 LIMIT 500000, will give top 50
0000 data points
# you can change the number to any other number based on your computing
power

# filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Sco
re != 3 LIMIT 500000""", con)
# for tsne assignment you can take 5k data points

filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score
!= 3 LIMIT 5000""", con)

# Give reviews with Score>3 a positive rating(1), and reviews with a sc
ore<3 a negative rating(0).
def partition(x):
    if x < 3:
        return 0
    return 1

```

```
#changing reviews with score less than 3 to be positive and vice-versa
actualScore = filtered_data['Score']
positiveNegative = actualScore.map(partition)
filtered_data['Score'] = positiveNegative
print("Number of data points in our data", filtered_data.shape)
filtered_data.head(3)
```

Number of data points in our data (5000, 10)

Out[0]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenomin
0	1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian	1	
1	2	B00813GRG4	A1D87F6ZCVE5NK	dll pa	0	
2	3	B000LQOCH0	ABXLMWJIXXAIN	Natalia Corres "Natalia Corres"	1	

```
In [0]: display = pd.read_sql_query("""
SELECT UserId, ProductId, ProfileName, Time, Score, Text, COUNT(*)
FROM Reviews
GROUP BY UserId
HAVING COUNT(*)>1
""", con)
```

```
In [0]: print(display.shape)
display.head()
```

(80668, 7)


Out[0]:

	UserId	ProductId	ProfileName	Time	Score	Text	COUNT(*)
0	#oc-R115TNMSPFT9I7	B007Y59HVM	Breyton	1331510400	2	Overall its just OK when considering the price...	2
1	#oc-R11D9D7SHXIJB9	B005HG9ET0	Louis E. Emory "hoppy"	1342396800	5	My wife has recurring extreme muscle spasms, u...	3
2	#oc-R11DNU2NBKQ23Z	B007Y59HVM	Kim Cieszykowski	1348531200	1	This coffee is horrible and unfortunately not ...	2
3	#oc-R11O5J5ZVQE25C	B005HG9ET0	Penguin Chick	1346889600	5	This will be the bottle that you grab from the...	3
4	#oc-R12KPBODL2B5ZD	B007OSBE1U	Christopher P. Presta	1348617600	1	I didnt like this coffee. Instead of telling y...	2

```
In [0]: display[display['UserId']=='AZY10LLTJ71NX']
```

Out[0]:

	UserId	ProductId	ProfileName	Time	Score	Text	COUNT(*)
80638	AZY10LLTJ71NX	B006P7E5ZI	undertheshrine "undertheshrine"	1334707200	5	I was recommended to try green tea extract to ...	5



```
In [0]: display['COUNT(*)'].sum()
```

Out[0]: 393063

[2] Exploratory Data Analysis

[2.1] Data Cleaning: Deduplication

It is observed (as shown in the table below) that the reviews data had many duplicate entries. Hence it was necessary to remove duplicates in order to get unbiased results for the analysis of the data. Following is an example:

```
In [0]: display= pd.read_sql_query("""  
SELECT *  
FROM Reviews  
WHERE Score != 3 AND UserId="AR5J8UI46CURR"  
ORDER BY ProductID  
""", con)  
display.head()
```

Out[0]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenon
0	78445	B000HDL1RQ	AR5J8UI46CURR	Geetha Krishnan	2	
1	138317	B000HDOPYC	AR5J8UI46CURR	Geetha Krishnan	2	
2	138277	B000HDOPYM	AR5J8UI46CURR	Geetha Krishnan	2	

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenon
3	73791	B000HDOPZG	AR5J8UI46CURR	Geetha Krishnan	2	
4	155049	B000PAQ75C	AR5J8UI46CURR	Geetha Krishnan	2	

As it can be seen above that same user has multiple reviews with same values for HelpfulnessNumerator, HelpfulnessDenominator, Score, Time, Summary and Text and on doing analysis it was found that

ProductId=B000HDOPZG was Loacker Quadratini Vanilla Wafer Cookies, 8.82-Ounce Packages (Pack of 8)

ProductId=B000HDL1RQ was Loacker Quadratini Lemon Wafer Cookies, 8.82-Ounce Packages (Pack of 8) and so on

It was inferred after analysis that reviews with same parameters other than ProductId belonged to the same product just having different flavour or quantity. Hence in order to reduce redundancy it was decided to eliminate the rows having same parameters.

The method used for the same was that we first sort the data according to ProductId and then just keep the first similar product review and delete the others. for eg. in the above just the review for ProductId=B000HDL1RQ remains. This method ensures that there is only one representative for each product and deduplication without sorting would lead to possibility of different representatives still existing for the same product.

```
In [0]: #Sorting data according to ProductId in ascending order
sorted_data=filtered_data.sort_values('ProductId', axis=0, ascending=True,
inplace=False, kind='quicksort', na_position='last')
```



```
In [0]: #Deduplication of entries
final=sorted_data.drop_duplicates(subset={"UserId","ProfileName","Time",
"Text"}, keep='first', inplace=False)
final.shape
```

Out[0]: (4986, 10)

```
In [0]: #Checking to see how much % of data still remains
(final['Id'].size*1.0)/(filtered_data['Id'].size*1.0)*100
```

Out[0]: 99.72

Observation:- It was also seen that in two rows given below the value of HelpfulnessNumerator is greater than HelpfulnessDenominator which is not practically possible hence these two rows too are removed from calculations

```
In [0]: display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND Id=44737 OR Id=64422
ORDER BY ProductID
""", con)

display.head()
```

Out[0]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenom
0	64422	B000MIDROQ	A161DK06JJMCYF	J. E. Stephens "Jeanne"	3	

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenomr
1	44737	B001EQ55RW	A2V0I904FH7ABY	Ram	3	

```
In [0]: final=final[final.HelpfulnessNumerator<=final.HelpfulnessDenominator]
```

```
In [0]: #Before starting the next phase of preprocessing lets see the number of
entries left
print(final.shape)

#How many positive and negative reviews are present in our dataset?
final['Score'].value_counts()
```

```
(4986, 10)
```

```
Out[0]: 1    4178
0     808
Name: Score, dtype: int64
```

[3] Preprocessing

[3.1]. Preprocessing Review Text

Now that we have finished deduplication our data requires some preprocessing before we go on further with analysis and making the prediction model.

Hence in the Preprocessing phase we do the following in the order below:-

1. Begin by removing the html tags
2. Remove any punctuations or limited set of special characters like , or . or # etc.

3. Check if the word is made up of english letters and is not alpha-numeric
4. Check to see if the length of the word is greater than 2 (as it was researched that there is no adjective in 2-letters)
5. Convert the word to lowercase
6. Remove Stopwords
7. Finally Snowball Stemming the word (it was observed to be better than Porter Stemming)

After which we collect the words used to describe positive and negative reviews

```
In [0]: # printing some random reviews
sent_0 = final['Text'].values[0]
print(sent_0)
print("="*50)

sent_1000 = final['Text'].values[1000]
print(sent_1000)
print("="*50)

sent_1500 = final['Text'].values[1500]
print(sent_1500)
print("="*50)

sent_4900 = final['Text'].values[4900]
print(sent_4900)
print("="*50)
```

Why is this \$[...] when the same product is available for \$[...] here?
<http://www.amazon.com/VICTOR-FLY-MAGNET-BAIT-REFILL/dp/B00004RBDY>

The Victor M380 and M502 traps are unreal, of course -- total fly genocide. Pretty stinky, but only right nearby.

=====

I recently tried this flavor/brand and was surprised at how delicious these chips are. The best thing was that there were a lot of "brown" chips in the bag (my favorite), so I bought some more through amazon and shared with family and friends. I am a little disappointed that there are not, so far, very many brown chips in these bags, but the flavor is still very good. I like them better than the yogurt and green onion flavor because they do not seem to be as salty, and the onion flavor is b

etter. If you haven't eaten Kettle chips before, I recommend that you try a bag before buying bulk. They are thicker and crunchier than Lays but just as fresh out of the bag.

=====

Wow. So far, two two-star reviews. One obviously had no idea what they were ordering; the other wants crispy cookies. Hey, I'm sorry; but these reviews do nobody any good beyond reminding us to look before ordering.

These are chocolate-oatmeal cookies. If you don't like that combination, don't order this type of cookie. I find the combo quite nice, really. The oatmeal sort of "calms" the rich chocolate flavor and gives the cookie sort of a coconut-type consistency. Now let's also remember that tastes differ; so, I've given my opinion.

Then, these are soft, chewy cookies -- as advertised. They are not "crispy" cookies, or the blurb would say "crispy," rather than "chewy." I happen to like raw cookie dough; however, I don't see where these taste like raw cookie dough. Both are soft, however, so is this the confusion? And, yes, they stick together. Soft cookies tend to do that. They aren't individually wrapped, which would add to the cost. Oh yeah, chocolate chip cookies tend to be somewhat sweet.

So, if you want something hard and crisp, I suggest Nabisco's Ginger Snaps. If you want a cookie that's soft, chewy and tastes like a combination of chocolate and oatmeal, give these a try. I'm here to place my second order.

=====

I love to order my coffee on amazon. easy and shows up quickly.
This k cup is great coffee. dcaf is very good as well

=====

```
In [0]: # remove urls from text python: https://stackoverflow.com/a/40823105/4084039
sent_0 = re.sub(r"http\S+", "", sent_0)
sent_1000 = re.sub(r"http\S+", "", sent_1000)
sent_150 = re.sub(r"http\S+", "", sent_150)
sent_4900 = re.sub(r"http\S+", "", sent_4900)

print(sent_0)
```

Why is this \$[...] when the same product is available for \$[...] here?

The Victor M380 and M502 traps are unreal, of course -- total fly genocide. Pretty stinky, but only right nearby.

```
In [0]: # https://stackoverflow.com/questions/16206380/python-beautifulsoup-how
        # -to-remove-all-tags-from-an-element
        from bs4 import BeautifulSoup

        soup = BeautifulSoup(sent_0, 'lxml')
        text = soup.get_text()
        print(text)
        print("="*50)

        soup = BeautifulSoup(sent_1000, 'lxml')
        text = soup.get_text()
        print(text)
        print("="*50)

        soup = BeautifulSoup(sent_1500, 'lxml')
        text = soup.get_text()
        print(text)
        print("="*50)

        soup = BeautifulSoup(sent_4900, 'lxml')
        text = soup.get_text()
        print(text)
```

Why is this \$[...] when the same product is available for \$[...] here?
 />The Victor M380 and M502 traps are unreal, of course -- total fly gen
 ocide. Pretty stinky, but only right nearby.

=====

I recently tried this flavor/brand and was surprised at how delicious t
 hese chips are. The best thing was that there were a lot of "brown" ch
 ips in the bsg (my favorite), so I bought some more through amazon and
 shared with family and friends. I am a little disappointed that there
 are not, so far, very many brown chips in these bags, but the flavor is
 still very good. I like them better than the yogurt and green onion fl
 avor because they do not seem to be as salty, and the onion flavor is b
 etter. If you haven't eaten Kettle chips before, I recommend that you
 try a bag before buying bulk. They are thicker and crunchier than Lays
 but just as fresh out of the bag.

=====

Wow. So far, two two-star reviews. One obviously had no idea what the

y were ordering; the other wants crispy cookies. Hey, I'm sorry; but these reviews do nobody any good beyond reminding us to look before ordering. These are chocolate-oatmeal cookies. If you don't like that combination, don't order this type of cookie. I find the combo quite nice, really. The oatmeal sort of "calms" the rich chocolate flavor and gives the cookie sort of a coconut-type consistency. Now let's also remember that tastes differ; so, I've given my opinion. Then, these are soft, chewy cookies -- as advertised. They are not "crispy" cookies, or the blurb would say "crispy," rather than "chewy." I happen to like raw cookie dough; however, I don't see where these taste like raw cookie dough. Both are soft, however, so is this the confusion? And, yes, they stick together. Soft cookies tend to do that. They aren't individually wrapped, which would add to the cost. Oh yeah, chocolate chip cookies tend to be somewhat sweet. So, if you want something hard and crisp, I suggest Nabisco's Ginger Snaps. If you want a cookie that's soft, chewy and tastes like a combination of chocolate and oatmeal, give these a try. I'm here to place my second order.

=====

love to order my coffee on amazon. easy and shows up quickly. This k cup is great coffee. dcaf is very good as well

```
In [0]: # https://stackoverflow.com/a/47091490/4084039
import re
```

```
def decontracted(phrase):
    # specific
    phrase = re.sub(r"won't", "will not", phrase)
    phrase = re.sub(r"can't", "can not", phrase)

    # general
    phrase = re.sub(r"n't", " not", phrase)
    phrase = re.sub(r"\ 're", " are", phrase)
    phrase = re.sub(r"\ 's", " is", phrase)
    phrase = re.sub(r"\ 'd", " would", phrase)
    phrase = re.sub(r"\ 'll", " will", phrase)
    phrase = re.sub(r"\ 't", " not", phrase)
    phrase = re.sub(r"\ 've", " have", phrase)
    phrase = re.sub(r"\ 'm", " am", phrase)
    return phrase
```

```
In [0]: sent_1500 = decontracted(sent_1500)
print(sent_1500)
print("="*50)
```

Wow. So far, two two-star reviews. One obviously had no idea what they were ordering; the other wants crispy cookies. Hey, I am sorry; but these reviews do nobody any good beyond reminding us to look before ordering.

These are chocolate-oatmeal cookies. If you do not like that combination, do not order this type of cookie. I find the combo quite nice, really. The oatmeal sort of "calms" the rich chocolate flavor and gives the cookie sort of a coconut-type consistency. Now let it also remember that tastes differ; so, I have given my opinion.

Then, these are soft, chewy cookies -- as advertised. They are not "crispy" cookies, or the blurb would say "crispy," rather than "chewy." I happen to like raw cookie dough; however, I do not see where these taste like raw cookie dough. Both are soft, however, so is this the confusion? And, yes, they stick together. Soft cookies tend to do that. They are not individually wrapped, which would add to the cost. Oh yeah, chocolate chip cookies tend to be somewhat sweet.

So, if you want something hard and crisp, I suggest Nabisco's Ginger Snaps. If you want a cookie that is soft, chewy and tastes like a combination of chocolate and oatmeal, give these a try. I am here to place my second order.

=====

```
In [0]: #remove words with numbers python: https://stackoverflow.com/a/18082370/4084039
sent_0 = re.sub("\S*\d\S*", "", sent_0).strip()
print(sent_0)
```

Why is this \$[...] when the same product is available for \$[...] here?

The Victor and traps are unreal, of course -- total fly genocide. Pretty stinky, but only right nearby.

```
In [0]: #remove spacial character: https://stackoverflow.com/a/5843547/4084039
sent_1500 = re.sub('[^A-Za-z0-9]+', ' ', sent_1500)
print(sent_1500)
```

Wow So far two two star reviews One obviously had no idea what they were ordering the other wants crispy cookies Hey I am sorry but these reviews do nobody any good beyond reminding us to look before ordering br br These are chocolate oatmeal cookies If you do not like that combination do not order this type of cookie I find the combo quite nice really The oatmeal sort of calms the rich chocolate flavor and gives the cookie sort of a coconut type consistency Now let us also remember that tastes differ so I have given my opinion br br Then these are soft chewy cookies as advertised They are not crispy cookies or the blurb would say crispy rather than chewy I happen to like raw cookie dough however I do not see where these taste like raw cookie dough Both are soft however so is this the confusion And yes they stick together Soft cookies tend to do that They are not individually wrapped which would add to the cost Oh yeah chocolate chip cookies tend to be somewhat sweet br br So if you want something hard and crisp I suggest Nabisco's Ginger Snaps If you want a cookie that is soft chewy and tastes like a combination of chocolate and oatmeal give these a try I am here to place my second order

```
In [0]: # https://gist.github.com/sebleier/554280
# we are removing the words from the stop words list: 'no', 'nor', 'not'
# <br /><br /> ==> after the above steps, we are getting "br br"
# we are including them into stop words list
# instead of <br /> if we have <br/> these tags would have been removed in the 1st step

stopwords= set(['br', 'the', 'i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves', 'you', "you're", "you've", \
               "you'll", "you'd", 'your', 'yours', 'yourself', 'yourselves', 'he', 'him', 'his', 'himself', \
               'she', "she's", 'her', 'hers', 'herself', 'it', "it's", 'its', 'itself', 'they', 'them', 'their', \
               'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'this', 'that', "that'll", 'these', 'those', \
               'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'have', 'has', 'had', 'having', 'do', 'does', \
               'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or', 'because', 'as', 'until', 'while', 'of', \
               'at', 'by', 'for', 'with', 'about', 'against', 'between',
```



```
'into', 'through', 'during', 'before', 'after', \
    'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out',
'on', 'off', 'over', 'under', 'again', 'further', \
    'then', 'once', 'here', 'there', 'when', 'where', 'why', 'h
ow', 'all', 'any', 'both', 'each', 'few', 'more', \
    'most', 'other', 'some', 'such', 'only', 'own', 'same', 's
o', 'than', 'too', 'very', \
    's', 't', 'can', 'will', 'just', 'don', "don't", 'should',
"should've", 'now', 'd', 'll', 'm', 'o', 're', \
    've', 'y', 'ain', 'aren', "aren't", 'couldn', "couldn't",
'didn', "didn't", 'doesn', "doesn't", 'hadn', \
    'hadn't', 'hasn', "hasn't", 'haven', "haven't", 'isn', "is
n't", 'ma', 'mightn', "mightn't", 'mustn', \
    "mustn't", 'needn', "needn't", 'shan', "shan't", 'shouldn',
"shouldn't", 'wasn', "wasn't", 'weren', "weren't", \
    'won', "won't", 'wouldn', "wouldn't"])
```

```
In [0]: # Combining all the above students
from tqdm import tqdm
preprocessed_reviews = []
# tqdm is for printing the status bar
for sentence in tqdm(final['Text'].values):
    sentence = re.sub(r"http\S+", "", sentence)
    sentence = BeautifulSoup(sentence, 'lxml').get_text()
    sentence = decontracted(sentence)
    sentence = re.sub("\S*\d\S*", "", sentence).strip()
    sentence = re.sub('[^A-Za-z]+', ' ', sentence)
    # https://gist.github.com/sebleier/554280
    sentence = ' '.join(e.lower() for e in sentence.split() if e.lower()
() not in stopwords)
    preprocessed_reviews.append(sentence.strip())
```

```
100%|██████████| 4986/4986 [00:01<00:00, 3137.37it/s]
```

```
In [0]: preprocessed_reviews[1500]
```

```
Out[0]: 'wow far two two star reviews one obviously no idea ordering wants cris
py cookies hey sorry reviews nobody good beyond reminding us look order'
```

```
ing chocolate oatmeal cookies not like combination not order type cooki
e find combo quite nice really oatmeal sort calms rich chocolate flavor
gives cookie sort coconut type consistency let also remember tastes dif
fer given opinion soft chewy cookies advertised not crispy cookies blur
b would say crispy rather chewy happen like raw cookie dough however no
t see taste like raw cookie dough soft however confusion yes stick toge
ther soft cookies tend not individually wrapped would add cost oh yeah
chocolate chip cookies tend somewhat sweet want something hard crisp su
ggest nabisco ginger snaps want cookie soft chewy tastes like combinatio
n chocolate oatmeal give try place second order'
```

[3.2] Preprocessing Review Summary

```
In [0]: ## Similarly you can do preprocessing for review summary also.
```

[4] Featurization

[4.1] BAG OF WORDS

```
In [0]: #BoW
count_vect = CountVectorizer() #in scikit-learn
count_vect.fit(preprocessed_reviews)
print("some feature names ", count_vect.get_feature_names()[:10])
print('='*50)

final_counts = count_vect.transform(preprocessed_reviews)
print("the type of count vectorizer ", type(final_counts))
print("the shape of out text BOW vectorizer ", final_counts.get_shape())
print("the number of unique words ", final_counts.get_shape()[1])

some feature names ['aa', 'aahhs', 'aback', 'abandon', 'abates', 'abb
ott', 'abby', 'abdominal', 'abiding', 'ability']
=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
```

```
the shape of out text BOW vectorizer (4986, 12997)
the number of unique words 12997
```

[4.2] Bi-Grams and n-Grams.

```
In [0]: #bi-gram, tri-gram and n-gram

#removing stop words like "not" should be avoided before building n-grams
# count_vect = CountVectorizer(ngram_range=(1,2))
# please do read the CountVectorizer documentation http://scikit-learn.org/stable/modules/generated/sklearn.feature\_extraction.text.CountVectorizer.html

# you can choose these numebrs min_df=10, max_features=5000, of your choice
count_vect = CountVectorizer(ngram_range=(1,2), min_df=10, max_features=5000)
final_bigram_counts = count_vect.fit_transform(preprocessed_reviews)
print("the type of count vectorizer ",type(final_bigram_counts))
print("the shape of out text BOW vectorizer ",final_bigram_counts.get_shape())
print("the number of unique words including both unigrams and bigrams ",
      final_bigram_counts.get_shape()[1])
```

```
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer (4986, 3144)
the number of unique words including both unigrams and bigrams 3144
```

[4.3] TF-IDF

```
In [0]: tf_idf_vect = TfidfVectorizer(ngram_range=(1,2), min_df=10)
tf_idf_vect.fit(preprocessed_reviews)
print("some sample features(unique words in the corpus)",tf_idf_vect.get_feature_names()[0:10])
print('='*50)
```

```

final_tf_idf = tf_idf_vect.transform(preprocessed_reviews)
print("the type of count vectorizer ",type(final_tf_idf))
print("the shape of out text TFIDF vectorizer ",final_tf_idf.get_shape
())
print("the number of unique words including both unigrams and bigrams "
, final_tf_idf.get_shape()[1])

```

some sample features(unique words in the corpus) ['ability', 'able', 'able find', 'able get', 'absolute', 'absolutely', 'absolutely delicious', 'absolutely love', 'absolutely no', 'according']

```

=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer (4986, 3144)
the number of unique words including both unigrams and bigrams 3144

```

[4.4] Word2Vec

```

In [0]: # Train your own Word2Vec model using your own text corpus
i=0
list_of_sentence=[]
for sentence in preprocessed_reviews:
    list_of_sentence.append(sentence.split())

```

```

In [0]: # Using Google News Word2Vectors

# in this project we are using a pretrained model by google
# its 3.3G file, once you load this into your memory
# it occupies ~9Gb, so please do this step only if you have >12G of ram
# we will provide a pickle file wich contains a dict ,
# and it contains all our courpus words as keys and model[word] as values
# To use this code-snippet, download "GoogleNews-vectors-negative300.bin"
# from https://drive.google.com/file/d/0B7XkCwpI5KDYNlNUTTlSS21pQmM/edit
# it's 1.9GB in size.

```

```

# http://kavita-ganesan.com/gensim-word2vec-tutorial-starter-code/#.W17
SRFAzZPY
# you can comment this whole cell
# or change these variable according to your need

is_your_ram_gt_16g=False
want_to_use_google_w2v = False
want_to_train_w2v = True

if want_to_train_w2v:
    # min_count = 5 considers only words that occurred at least 5 times
    w2v_model=Word2Vec(list_of_sentence,min_count=5,size=50, workers=4)
    print(w2v_model.wv.most_similar('great'))
    print('='*50)
    print(w2v_model.wv.most_similar('worst'))

elif want_to_use_google_w2v and is_your_ram_gt_16g:
    if os.path.isfile('GoogleNews-vectors-negative300.bin'):
        w2v_model=KeyedVectors.load_word2vec_format('GoogleNews-vectors
-negative300.bin', binary=True)
        print(w2v_model.wv.most_similar('great'))
        print(w2v_model.wv.most_similar('worst'))
    else:
        print("you don't have google's word2vec file, keep want_to_train_w2v = True, to train your own w2v ")

[('snack', 0.9951335191726685), ('calorie', 0.9946465492248535), ('wonderful', 0.9946032166481018), ('excellent', 0.9944332838058472), ('especially', 0.9941144585609436), ('baked', 0.9940600395202637), ('salted', 0.994047224521637), ('alternative', 0.9937226176261902), ('tasty', 0.9936816692352295), ('healthy', 0.9936649799346924)]
=====
[('varieties', 0.9994194507598877), ('become', 0.9992934465408325), ('popcorn', 0.9992750883102417), ('de', 0.9992610216140747), ('miss', 0.9992451071739197), ('melitta', 0.999218761920929), ('choice', 0.9992102384567261), ('american', 0.9991837739944458), ('beef', 0.9991780519485474), ('finish', 0.9991567134857178)]

```

```
In [0]: w2v_words = list(w2v_model.wv.vocab)
print("number of words that occurred minimum 5 times ", len(w2v_words))
print("sample words ", w2v_words[0:50])
```

number of words that occurred minimum 5 times 3817
sample words ['product', 'available', 'course', 'total', 'pretty', 'st
inky', 'right', 'nearby', 'used', 'ca', 'not', 'beat', 'great', 'receiv
ed', 'shipment', 'could', 'hardly', 'wait', 'try', 'love', 'call', 'ins
tead', 'removed', 'easily', 'daughter', 'designed', 'printed', 'use',
'car', 'windows', 'beautifully', 'shop', 'program', 'going', 'lot', 'fu
n', 'everywhere', 'like', 'tv', 'computer', 'really', 'good', 'idea',
'final', 'outstanding', 'window', 'everybody', 'asks', 'bought', 'mad
e']

[4.4.1] Converting text into vectors using Avg W2V, TFIDF-W2V

[4.4.1.1] Avg W2v

```
In [0]: # average Word2Vec
# compute average word2vec for each review.
sent_vectors = []; # the avg-w2v for each sentence/review is stored in
this list
for sent in tqdm(list_of_sentence): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length 50, yo
u might need to change this to 300 if you use google's w2v
    cnt_words = 0; # num of words with a valid vector in the sentence/re
view
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    sent_vectors.append(sent_vec)
```

100% | ██
██████████ | 4986/4986 [00:03<00:00, 1330.47it/s]

[4.4.1.2] TFIDF weighted W2v

```
In [0]: # TF-IDF weighted Word2Vec
tfidf_feat = model.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and ce
ll_val = tfidf

tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review is st
ored in this list
row=0;
for sent in tqdm(list_of_sentence): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/r
review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
#             tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
# to reduce the computation we are
# dictionary[word] = idf value of word in whole courpus
# sent.count(word) = tf valeus of word in this review
            tf_idf = dictionary[word]*(sent.count(word)/len(sent))
```

```
100%|██████████| 4986/4986 [00:20<00:00, 245.63it/s]
```


- Add a small eps value(to eliminate the divisible by zero error) to W and W' i.e $W=W+10^{-6}$ and $W' = W'+10^{-6}$
- Now find the % change between W and W' ($(W-W') / (W) \times 100$)
- Calculate the 0th, 10th, 20th, 30th, ...100th percentiles, and observe any sudden rise in the values of percentage_change_vector
- Ex: consider your 99th percentile is 1.3 and your 100th percentiles are 34.6, there is sudden rise from 1.3 to 34.6, now calculate the 99.1, 99.2, 99.3,..., 100th percentile values and get the proper value after which there is sudden rise the values, assume it is 2.5
- Print the feature names whose % change is more than a threshold x(in our example it's 2.5)

4. Sparsity

- Calculate sparsity on weight vector obtained after using L1 regularization

NOTE: Do sparsity and multicollinearity for any one of the vectorizers. Bow or tf-idf is recommended.

5. Feature importance

- Get top 10 important features for both positive and negative classes separately.

6. Feature engineering


- To increase the performance of your model, you can also experiment with with feature engineering like :
 - Taking length of reviews as another feature.
 - Considering some features from review summary as well.

7. Representation of results

- You need to plot the performance of model both on train data and cross validation data for each hyper parameter, like shown in the figure.



Once after you found the best hyper parameter, you need to train your model with it, and find the AUC on test data and plot the ROC curve on both train and test.

 Along with plotting ROC curve, you need to print the [confusion matrix](#) with predicted and original labels of test data points. Please visualize your confusion matrices using [seaborn heatmaps](#).



8. [Conclusion](#)

- [You need to summarize the results at the end of the notebook, summarize it in the table format. To print out a table please refer to this prettytable library link](#)



Note: Data Leakage

1. There will be an issue of data-leakage if you vectorize the entire data and then split it into train/cv/test.
2. To avoid the issue of data-leakag, make sure to split your data first and then vectorize it.
3. While vectorizing your data, apply the method `fit_transform()` on you train data, and apply the method `transform()` on cv/test data.
4. For more details please go through this [link](#).

Applying Logistic Regression

[5.1] Logistic Regression on BOW, SET 1

[5.1.1] Applying Logistic Regression with L1 regularization on BOW, SET 1

In [0]: `final.columns`

```
Out[0]: Index(['index', 'Id', 'ProductId', 'UserId', 'ProfileName',
              'HelpfulnessNumerator', 'HelpfulnessDenominator', 'Score', 'Tim
              e',
              'Summary', 'Text', 'CleanedText'],
              dtype='object')
```

```
In [0]: from sklearn.model_selection import GridSearchCV
        from sklearn.model_selection import RandomizedSearchCV
        from sklearn.model_selection import TimeSeriesSplit
        from sklearn.model_selection import train_test_split

        preprocessed_reviews=final['CleanedText'][:100000]
        score=final['Score'][:100000]
        X_train, X_test, y_train, y_test = train_test_split(preprocessed_reviews,
        score, test_size=0.25,stratify=score, random_state=42)
```

```
In [8]: #Bow
        count_vect = CountVectorizer(max_df=0.95, min_df=2,stop_words='english',
        ,max_features=10000) #in scikit-learn
        count_vect.fit(X_train)
        print("some feature names ", count_vect.get_feature_names()[:10])
        print('='*50)

        X_train_bow = count_vect.transform(X_train)
        print("the type of count vectorizer ",type(X_train_bow))
        print("the shape of out text BOW vectorizer ",X_train_bow.get_shape())
        print("the number of unique words ", X_train_bow.get_shape()[1])

        X_test_bow = count_vect.transform(X_test)
        print("the type of count vectorizer ",type(X_test_bow))
        print("the shape of out text BOW vectorizer ",X_test_bow.get_shape())
        print("the number of unique words ", X_test_bow.get_shape()[1])

        some feature names  ['ab', 'aback', 'abandon', 'abbey', 'abc', 'abdomi
        n', 'abil', 'abl', 'abnorm', 'abomin']
        =====
        the type of count vectorizer  <class 'scipy.sparse.csr.csr_matrix'>
        the shape of out text BOW vectorizer  (75000, 10000)
        the number of unique words  10000
```

```
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer (25000, 10000)
the number of unique words 10000
```

```
In [0]: #X_train_bow=X_train_bow.toarray()

from sklearn.preprocessing import StandardScaler
scaler = StandardScaler(with_mean=False)
scaler.fit(X_train_bow)
X_train_bow=scaler.transform(X_train_bow)

X_test_bow=scaler.transform(X_test_bow)
```

```
In [0]: from sklearn.linear_model import LogisticRegression

clf=LogisticRegression(penalty='l1')
param_grid={'C':[1000,100,10,1,0.1,0.001,0.0001]}
#timeseriessplit=TimeSeriesSplit(n_splits=10)
gcv=GridSearchCV(clf,param_grid,cv=5,scoring='roc_auc')
gcv.fit(X_train_bow,y_train)
print(gcv.best_params_)
print(gcv.best_score_)
```

```
In [0]: hyper_parameters=gcv.get_params()['param_grid']['C']
train_scores=gcv.cv_results_['mean_train_score'].tolist()
test_scores=gcv.cv_results_['mean_test_score'].tolist()

print(hyper_parameters)
print(test_scores)
print(train_scores)

[1000, 100, 10, 5, 1, 0.5, 0.1, 0.001, 0.0001]
[0.9088546452748788, 0.9103610412514892, 0.912518280012368, 0.913361678
1239468, 0.9157905093369368, 0.9173113487846236, 0.9236798904189109, 0.
9129397282067238, 0.6950984355652179]
[0.9659050271342888, 0.9659028612935927, 0.9658948545387306, 0.96588897
56573034, 0.9658521776960267, 0.9658079422705249, 0.965298273599571, 0.
914513249597746, 0.695107385548616]
```

```

In [7]: import math
C = [1000, 100, 10, 5, 1, 0.5, 0.1, 0.001, 0.0001]
C_log=[math.log(i) for i in C]

test_scores=[0.9088546452748788, 0.9103610412514892, 0.912518280012368,
0.9133616781239468, 0.9157905093369368, 0.9173113487846236, 0.92367989
04189109, 0.9129397282067238, 0.6950984355652179]
train_scores=[0.9659050271342888, 0.9659028612935927, 0.965894854538730
6, 0.9658889756573034, 0.9658521776960267, 0.9658079422705249, 0.965298
273599571, 0.914513249597746, 0.695107385548616]

fig, ax = plt.subplots()
ax.plot(C_log, train_scores,c='g',marker='o',label="TrainAccuracy")

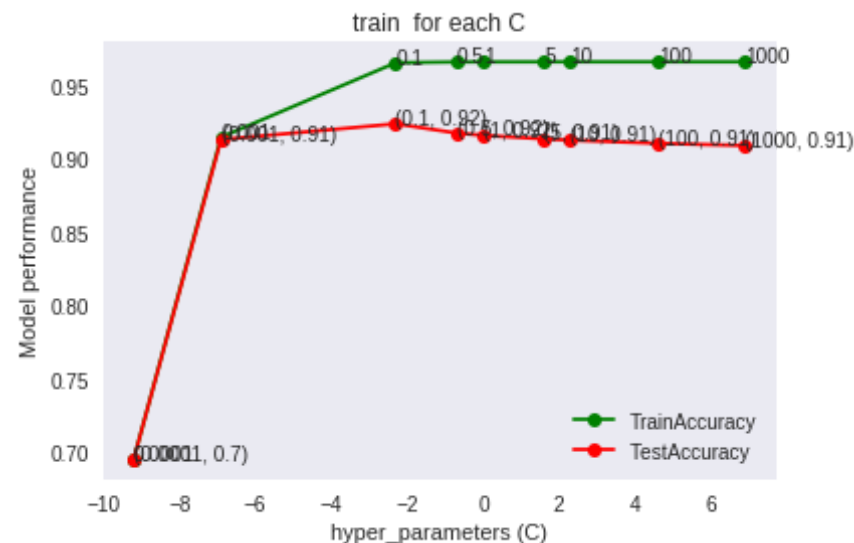
for i, txt in enumerate(C):
    ax.annotate(txt, (C_log[i], train_scores[i]))

ax.plot(C_log, test_scores ,c='r',marker='o',label="TestAccuracy")

for i, txt in enumerate(C):
    ax.annotate((txt,np.round(test_scores[i],2)) , (C_log[i], test_scores[i]))

plt.title("train for each C")
plt.xlabel("hyper_parameters (C)")
plt.ylabel("Model performance")
plt.legend()
plt.grid()
plt.show()

```



```
In [0]: from sklearn.metrics import roc_auc_score
from sklearn.metrics import auc
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
from sklearn.metrics import f1_score

clf1=LogisticRegression(C=0.1,penalty='l1')
clf1.fit(X_train_bow,y_train)
pred_train=clf1.predict(X_train_bow)
pred=clf1.predict(X_test_bow)

print("Accuracy Score : ",accuracy_score(y_test,pred)*100)
print("Precision Score : ",precision_score(y_test,pred)*100)
print("Recall Score : ",recall_score(y_test,pred)*100)
print("F1 Score : ",f1_score(y_test,pred)*100)

print("      ")
```

```

print("Classification Report")
print(classification_report(y_test,pred))
print(" ")

fpr_train,tpr_train,thresholds_train=roc_curve(y_train,pred_train)
print("AUC Score for train data :",metrics.auc(fpr_train,tpr_train))

fpr,tpr,thresholds=roc_curve(y_test,pred)
print("AUC Score for test data :",metrics.auc(fpr,tpr))

print(" ")

#y_true = # ground truth labels
#y_probas = # predicted probabilities generated by sklearn classifier
#skplt.metrics.plot_roc_curve(y_true, y_probas)
#plt.show()

plt.figure()
lw = 2
plt.plot(fpr, tpr, color='red',
         lw=lw,label='test')
plt.plot(fpr_train, tpr_train, color='darkorange',
         lw=lw,label='train')
plt.plot([0, 1], [0, 1], color='navy', lw=lw, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic example')
plt.legend(loc="lower right")
plt.show()

print(" ")

tn, fp, fn, tp=confusion_matrix(y_test,pred).ravel()
print("""

```

```

TrueNegative : {}
FalsePositive : {}
FalseNegative : {}
TruePositive : {}"".format(tn, fp, fn, tp))
print(" ")
print(" ")

confusionmatrix_DF=pd.DataFrame(confusion_matrix(y_test,pred),columns=[
'0','1'],index=['0','1'])
sns.heatmap(confusionmatrix_DF,annot=True,fmt='g',cmap='viridis')
plt.title("Confusion matrix ")
plt.show()


```

Accuracy Score : 90.708
 Precision Score : 93.72350230414746
 Recall Score : 95.48805108221043
 F1 Score : 94.59754878020419

Classification Report					
	precision	recall	f1-score	support	
0	0.71	0.63	0.67	3701	
1	0.94	0.95	0.95	21299	
micro avg	0.91	0.91	0.91	25000	
macro avg	0.82	0.79	0.81	25000	
weighted avg	0.90	0.91	0.90	25000	


AUC Score for train data : 0.8962604569087288
 AUC Score for test data : 0.7934359322551484

Receiver operating characteristic example



TrueNegative : 2339
FalsePositive : 1362
FalseNegative : 961
TruePositive : 20338





with PREROCESSED REVIEWS

```
In [0]: hyper_parameters=gcv.get_params()['param_grid']['C']
train_scores=gcv.cv_results_['mean_train_score'].tolist()
test_scores=gcv.cv_results_['mean_test_score'].tolist()


print(hyper_parameters)
print(test_scores)
print(train_scores)

[1000, 100, 10, 5, 1, 0.5, 0.1, 0.001, 0.0001]
[0.7908377521433998, 0.8079856879597342, 0.8181680236380514, 0.82162558
65348513, 0.8335232007342616, 0.8431101705158265, 0.8819334254079269,
0.5, 0.5]
[1.0, 1.0, 1.0, 1.0, 0.9999706414160041, 0.9997050308702843, 0.98806266
76516027, 0.5, 0.5]
```

```
In [0]: plt.plot( hyper_parameters ,train_scores , label='Train plot')
plt.plot( hyper_parameters ,test_scores , label='Test plot')
plt.xlabel("hyper_parameters (C)")
plt.ylabel("Model performance")

plt.legend()
```

Out[0]: <matplotlib.legend.Legend at 0x7ff8cc72bb00>



In [8]: `import math`

```
C = [1000, 100, 10, 5, 1, 0.5, 0.1, 0.001, 0.0001]
C_log=[math.log(i) for i in C]
test_scores=[0.7908377521433998, 0.8079856879597342, 0.8181680236380514,
, 0.8216255865348513, 0.8335232007342616, 0.8431101705158265, 0.8819334
254079269, 0.5, 0.5]
train_scores = [1.0, 1.0, 1.0, 1.0, 0.9999706414160041, 0.9997050308702
843, 0.9880626676516027, 0.5, 0.5]

fig, ax = plt.subplots()
ax.plot(C_log, train_scores,c='g',marker='o',label="TrainAccuracy")

for i, txt in enumerate(C):
    ax.annotate(txt, (C_log[i], train_scores[i]))

ax.plot(C_log, test_scores ,c='r',marker='o',label="TestAccuracy")

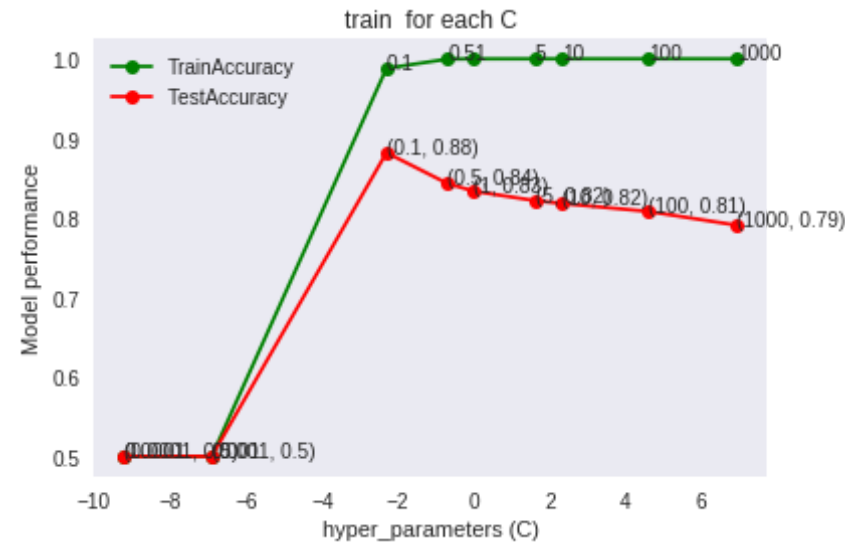
for i, txt in enumerate(C):
```

```

ax.annotate((txt,np.round(test_scores[i],2)) , (C_log[i], test_scores[i]))

plt.title("train for each C")
plt.xlabel("hyper_parameters (C)")
plt.ylabel("Model performance")
plt.legend()
plt.grid()
plt.show()

```



```

In [0]: from sklearn.metrics import roc_auc_score
        from sklearn.metrics import auc
        from sklearn.metrics import accuracy_score
        from sklearn.metrics import confusion_matrix
        from sklearn.metrics import classification_report
        from sklearn.metrics import precision_score
        from sklearn.metrics import recall_score
        from sklearn.metrics import f1_score

        clf1=LogisticRegression(C=0.1,penalty='l1')

```

```

clf1.fit(X_train_bow,y_train)
pred_train=clf1.predict(X_train_bow)
pred=clf1.predict(X_test_bow)

print("Accuracy Score : ",accuracy_score(y_test,pred)*100)
print("Precision Score : ",precision_score(y_test,pred)*100)
print("Recall Score : ",recall_score(y_test,pred)*100)
print("F1 Score : ",f1_score(y_test,pred)*100)

print(" ")
print("Classification Report")
print(classification_report(y_test,pred))
print(" ")

fpr_train,tpr_train,thresholds_train=roc_curve(y_train,pred_train)
print("AUC Score for train data :",metrics.auc(fpr_train,tpr_train))

fpr,tpr,thresholds=roc_curve(y_test,pred)
print("AUC Score for test data :",metrics.auc(fpr,tpr))

print(" ")

#y_true = # ground truth labels
#y_probas = # predicted probabilities generated by sklearn classifier
#skplt.metrics.plot_roc_curve(y_true, y_probas)
#plt.show()

plt.figure()
lw = 2
plt.plot(fpr, tpr, color='red',
         lw=lw,label='test')
plt.plot(fpr_train, tpr_train, color='darkorange',
         lw=lw,label='train')
plt.plot([0, 1], [0, 1], color='navy', lw=lw, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')

```

```

plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic example')
plt.legend(loc="lower right")
plt.show()

print(" ")

tn, fp, fn, tp=confusion_matrix(y_test,pred).ravel()
print("""
TrueNegative : {}
FalsePositive : {}
FalseNegative : {}
TruePositive : {}""".format(tn, fp, fn, tp))
print(" ")
print(" ")

confusionmatrix_DF=pd.DataFrame(confusion_matrix(y_test,pred),columns=[
'0','1'],index=['0','1'])
sns.heatmap(confusionmatrix_DF,annot=True,fmt='g',cmap='viridis')
plt.title("Confusion matrix ")
plt.show()


```

Accuracy Score : 88.0
Precision Score : 91.07883817427386
Recall Score : 95.02164502164501
F1 Score : 93.00847457627118


Classification Report				
	precision	recall	f1-score	support
0	0.66	0.51	0.58	528
1	0.91	0.95	0.93	2772
micro avg	0.88	0.88	0.88	3300
macro avg	0.79	0.73	0.75	3300
weighted avg	0.87	0.88	0.87	3300

AUC Score for train data : 0.8778612876448715

AUC Score for test data : 0.7307900432900433



TrueNegative : 270
FalsePositive : 258
FalseNegative : 138
TruePositive : 2634



[5.1.1.1] Calculating sparsity on weight vector obtained using L1 regularization on BOW, SET 1

```
In [0]: a=np.count_nonzero(clf1.coef_)
size=clf1.coef_.size
Sparsity= (size - a)/size
```

```
In [0]: hyper_parameters=gcv.get_params()['param_grid']['C']
train_scores=gcv.cv_results_['mean_train_score'].tolist()
test_scores=gcv.cv_results_['mean_test_score'].tolist()

print(hyper_parameters)
print(test_scores)
print(train_scores)

[1000, 100, 10, 5, 1, 0.5, 0.1, 0.001, 0.0001]
[0.9088546452748788, 0.9103610412514892, 0.912518280012368, 0.913361678
1239468, 0.9157905093369368, 0.9173113487846236, 0.9236798904189109, 0.
9129397282067238, 0.6950984355652179]
[0.9659050271342888, 0.9659028612935927, 0.9658948545387306, 0.96588897
```



```
56573034, 0.9658521776960267, 0.9658079422705249, 0.965298273599571, 0.914513249597746, 0.695107385548616]
```

In [0]:

[5.1.2] Applying Logistic Regression with L2 regularization on BOW, SET 1

In [0]: `from sklearn.linear_model import LogisticRegression`

```
clf=LogisticRegression(penalty='l2')
param_grid={'C':[1000,100,10,1,0.1,0.001,0.0001]}
#timeseriessplit=TimeSeriesSplit(n_splits=10)
gcv=GridSearchCV(clf,param_grid,cv=5,scoring='roc_auc')
gcv.fit(X_train_bow,y_train)
print(gcv.best_params_)
print(gcv.best_score_)
```

```
{'C': 0.001}
0.9097617785177461
```

In [0]: `hyper_parameters=gcv.get_params()['param_grid']['C']`
`train_scores=gcv.cv_results_['mean_train_score'].tolist()`
`test_scores=gcv.cv_results_['mean_test_score'].tolist()`


```
print(hyper_parameters)
print(test_scores)
print(train_scores)
```

```
[1000, 100, 10, 1, 0.1, 0.001, 0.0001]
[0.8511000043019294, 0.8506357843347192, 0.8536634128701011, 0.85332302
13012999, 0.8600912167589186, 0.9097617785177461, 0.90257602348846]
[0.9921086446080356, 0.992160164672204, 0.9920741278201515, 0.992154196
3998797, 0.9918990281685094, 0.9822747157003355, 0.9562920434285788]
```

In [0]: `plt.plot(hyper_parameters ,train_scores , label='Train plot')`
`plt.plot(hyper_parameters ,test_scores , label='Test plot')`
`plt.xlabel("hyper_parameters (C)")`

```
plt.ylabel("Model performance")
plt.legend()
```

Out[0]: <matplotlib.legend.Legend at 0x7fe9a9bc8438>



```
In [0]: from sklearn.metrics import roc_auc_score
from sklearn.metrics import auc
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
from sklearn.metrics import f1_score

clf1=LogisticRegression(C=0.001,penalty='l2')
clf1.fit(X_train_bow,y_train)
pred_train=clf1.predict(X_train_bow)
pred=clf1.predict(X_test_bow)

print("Accuracy Score : ",accuracy_score(y_test,pred)*100)
```

```

print("Precision Score : ",precision_score(y_test,pred)*100)
print("Recall Score : ",recall_score(y_test,pred)*100)
print("F1 Score : ",f1_score(y_test,pred)*100)

print(" ")
print("Classification Report")
print(classification_report(y_test,pred))
print(" ")

fpr_train,tpr_train,thresholds_train=roc_curve(y_train,pred_train)
print("AUC Score for train data :",metrics.auc(fpr_train,tpr_train))

fpr,tpr,thresholds=roc_curve(y_test,pred)
print("AUC Score for test data :",metrics.auc(fpr,tpr))

print(" ")

#y_true = # ground truth labels
#y_probas = # predicted probabilities generated by sklearn classifier
#skplt.metrics.plot_roc_curve(y_true, y_probas)
#plt.show()

plt.figure()
lw = 2
plt.plot(fpr, tpr, color='red',
         lw=lw,label='test')
plt.plot(fpr_train, tpr_train, color='darkorange',
         lw=lw,label='train')
plt.plot([0, 1], [0, 1], color='navy', lw=lw, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic example')
plt.legend(loc="lower right")
plt.show()

```

```

print("      ")

tn, fp, fn, tp=confusion_matrix(y_test,pred).ravel()
print("""
TrueNegative : {}
FalsePositive : {}
FalseNegative : {}
TruePositive : {}""".format(tn, fp, fn, tp))
print("      ")
print("      ")


confusionmatrix_DF=pd.DataFrame(confusion_matrix(y_test,pred),columns=[
'0','1'],index=['0','1'])
sns.heatmap(confusionmatrix_DF,annot=True,fmt='g',cmap='viridis')
plt.title("Confusion matrix ")
plt.show()

```

Accuracy Score : 91.268
Precision Score : 93.17463185472943
Recall Score : 96.84492229682145
F1 Score : 94.97433063977715

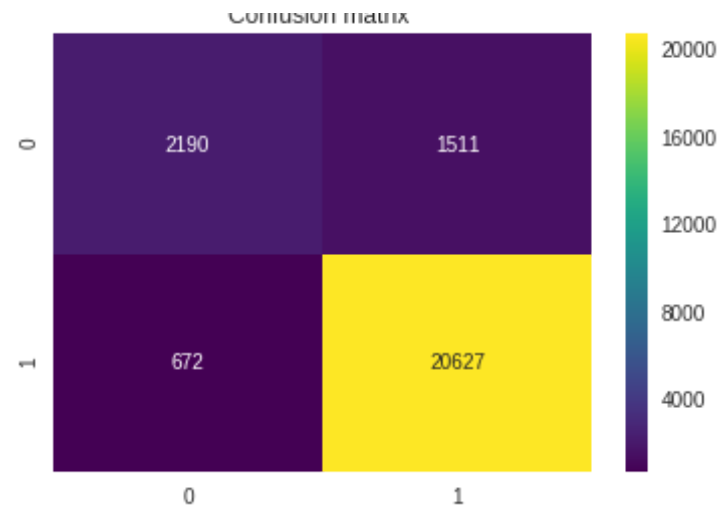
Classification Report				
	precision	recall	f1-score	support
0	0.77	0.59	0.67	3701
1	0.93	0.97	0.95	21299
micro avg	0.91	0.91	0.91	25000
macro avg	0.85	0.78	0.81	25000
weighted avg	0.91	0.91	0.91	25000

AUC Score for train data : 0.8652155813496196
AUC Score for test data : 0.7800905936510891



TrueNegative : 2190
FalsePositive : 1511
FalseNegative : 672
TruePositive : 20627

Confusion matrix



[5.1.2.1] Performing pertubation test (multicollinearity check) on BOW, SET 1

this noise can also be used but this gives dense matrix

```
In [0]: noise=0.000001
x_train_bow_noise=np.zeros(X_train_bow.shape)
for i in range(X_train_bow.shape[0]):
    x_train_bow_noise[i]=X_train_bow[i].toarray()+noise
```

```
In [12]: #x_train_bow_noise
from sklearn.linear_model import LogisticRegression
clfNO=LogisticRegression(C=0.001,penalty='l2')
clfNO.fit(X_train_bow,y_train)
print("WEIGHTS with X_train_bow :", clfNO.coef_[0])

WEIGHTS with X_train_bow : [ 0.00019553 -0.00740731 -0.00017604 ...  0.
0.00757321  0.0008448
0.00331951]
```

```
In [13]: clfY=LogisticRegression(C=0.001,penalty='l2')
```

```
clfY.fit(x_train_noisel,y_train)
print("WEIGHTS with x_train_bow_noise :",clfY.coef_[0])
```

```
WEIGHTS with x_train_bow_noise : [ 0.00019553 -0.00740731 -0.00017604
... 0.00757321 0.0008448
0.00331951]
```

I have used this noise

```
In [0]: import copy
x_train_noisel=copy.deepcopy(X_train_bow)
#e=np.random.normal(0,0.01)
x_train_noisel.data += 0.00001
```

```
In [93]: type(x_train_noisel)
```

```
Out[93]: scipy.sparse.csr.csr_matrix
```

```
In [97]: #weight without adding noise to data :::: so used X_train_bow
clfN0=LogisticRegression(C=0.001,penalty='l2')
clfN0.fit(X_train_bow,y_train)
print("weight without noise using X_train_bow : \n",clfN0.coef_[0])
print('\n')

#weight with adding noise to data :::: so used x_train_noisel
clfY=LogisticRegression(C=0.001,penalty='l2')
clfY.fit(x_train_noisel,y_train)
print("weight with noise using x_train_noisel :\n ",clfY.coef_[0])
print('\n')

#Calculated WeightDifference
wgts_difference=(abs((clfN0.coef_[0]-clfY.coef_[0])/clfN0.coef_[0])*100
)
print("weight Difference between noise and not noise :\n ",wgts_differe
nce)
print('\n')
print("counting any weight difference greater than 0.5 : ",wgts_differ
```

```
ence[np.where(wgts_difference > 0.5)].size)
print('\n')
print("length of Weight difference is : ",len(wgts_difference))
print('\n')
print("Weight difference is : \n",wgts_difference)
```

```
weight without noise using X_train_bow :
[ 0.00019553 -0.00740731 -0.00017604 ...  0.00757321  0.0008448
 0.00331951]
```

```
weight with noise using x_train_noisel :
[ 0.00019562 -0.00740731 -0.00017603 ...  0.00757322  0.00084479
 0.00331946]
```

```
weight Difference between noise and not noise :
[4.87617935e-02 2.08552250e-05 4.87784215e-03 ... 1.49830512e-04
 1.28442956e-03 1.58088976e-03]
```

```
counting any weight difference greater than 0.5 : 15
```

```
length of Weight difference is : 10000
```

```
Weight difference is :
[4.87617935e-02 2.08552250e-05 4.87784215e-03 ... 1.49830512e-04
 1.28442956e-03 1.58088976e-03]
```

```
In [98]: from sklearn.linear_model import LogisticRegression

#weight without adding noise to data :::: so used x_train_noisel
clf1=LogisticRegression(C=0.001,penalty='l2')
clf1.fit(X_train_bow,y_train)
pred=clf1.predict(X_test_bow)
a=np.count_nonzero(clf1.coef_)
print("AccuracyScore : ",accuracy_score(y_test,pred))
```



```

print(a)

#After getting weights adding 10**-6 to weights
#adding W=W+10^-6
print(clf1.coef_.shape)
wgts_without_noise=clf1.coef_[0]
print(wgts_without_noise)
wgts_without_noise=wgts_without_noise+0.000001
print(wgts_without_noise)

#x_train_noise1=copy.deepcopy(X_train_bow)
#x_train_noise1.data += 0.00001
#weight with adding noise to data :::: so used x_train_noise1
clf1_noise=LogisticRegression(C=0.001,penalty='l2')
clf1_noise.fit(x_train_noise1,y_train)
pred=clf1_noise.predict(X_test_bow)
a=np.count_nonzero(clf1_noise.coef_)
print("AccuracyScore : ",accuracy_score(y_test,pred))
print(a)

#After getting weights adding 10**-6 to weights
# adding W' = W'+10^-6
print(clf1_noise.coef_.shape)
wgts_with_noise=clf1_noise.coef_[0]
print(wgts_with_noise)
wgts_with_noise=wgts_with_noise+0.000001
print(wgts_with_noise)

wgts_difference=(abs((wgts_without_noise-wgts_with_noise)/wgts_without_noise)*100)
print(wgts_difference[np.where(wgts_difference > 0.5)].size)
print("Weight differeemce is : ",wgts_difference)

#Calculated WeightDifference
wgts_difference=(abs((wgts_without_noise-wgts_with_noise)/wgts_without_

```

```

noise)*100)
print("weight Difference between noise and not noise :\n ",wgts_difference)
print('\n')
print("counting any weight difference greater than 0.5 : ",wgts_difference[np.where(wgts_difference > 0.5)].size)
print('\n')
print("length of Weight difference is : ",len(wgts_difference))
print('\n')
print("Weight difference is : \n",wgts_difference)

```

AccuracyScore : 0.91268

10000

(1, 10000)

[0.00019553 -0.00740731 -0.00017604 ... 0.00757321 0.0008448
0.00331951]

[0.00019653 -0.00740631 -0.00017504 ... 0.00757421 0.0008458
0.00332051]

AccuracyScore : 0.91268

10000

(1, 10000)

[0.00019562 -0.00740731 -0.00017603 ... 0.00757322 0.00084479
0.00331946]

[0.00019662 -0.00740631 -0.00017503 ... 0.00757422 0.00084579
0.00332046]

7

Weight difference is : [4.85136782e-02 2.08580409e-05 4.90570968e-03
... 1.49810731e-04

1.28291097e-03 1.58041367e-03]

weight Difference between noise and not noise :

[4.85136782e-02 2.08580409e-05 4.90570968e-03 ... 1.49810731e-04
1.28291097e-03 1.58041367e-03]

counting any weight difference greater than 0.5 : 7

length of Weight difference is : 10000

```
Weight difference is :  
[4.85136782e-02 2.08580409e-05 4.90570968e-03 ... 1.49810731e-04  
1.28291097e-03 1.58041367e-03]
```

```
In [103]: #tag_df_sorted = tag_df.sort_values(['Counts'], ascending=False)  
#tag_counts = tag_df_sorted['Counts'].values
```

```
wgts_difference.sort()  
print(wgts_difference)
```

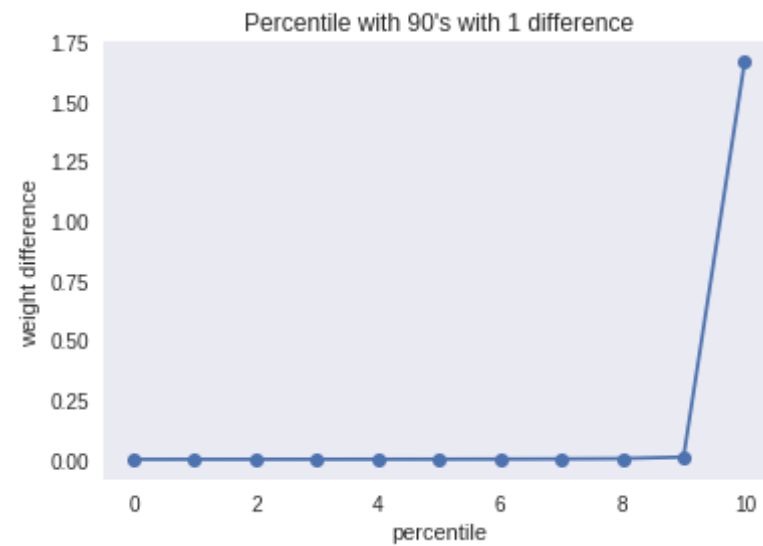

```
[2.95507520e-08 5.96602930e-08 7.91597000e-08 ... 9.64262480e-01  
1.43812690e+00 1.66461003e+00]
```

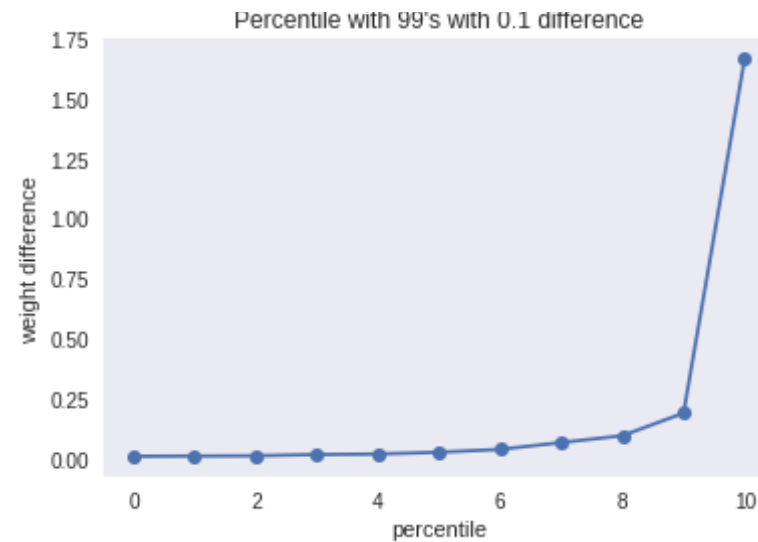
```
In [106]: #wgts_difference  
percentile10=np.percentile(wgts_difference, [0,10,20,30,40,50,60,70,80,  
90,100])  
plt.plot(percentile10,'-o')  
plt.title("Percentile with 10 difference")  
plt.grid()  
plt.xlabel("percentile")  
plt.ylabel("weight difference")  
plt.show()
```

```
b=np.percentile(wgts_difference,[90,91,92,93,94,95,96,97,98,99,100])  
plt.plot(b,'-o')  
plt.title("Percentile with 90's with 1 difference")  
plt.grid()  
plt.xlabel("percentile")  
plt.ylabel("weight difference")  
plt.show()
```

```
a=np.percentile(wgts_difference,[99,99.1,99.2,99.3,99.4,99.5,99.6,99.7,  
99.8,99.9,100])  
plt.plot(a,'-o')  
plt.title("Percentile with 99's with 0.1 difference")  
plt.grid()  
plt.xlabel("percentile")
```

```
plt.ylabel("weight difference")  
plt.show()
```





By the above graph sudden rise at 99.9 percentile

Another way of plotting percentiles

```
In [107]: plt.plot(wgts_difference, '-o')
plt.title("Percentile's plotting")
plt.grid()
plt.xlabel("percentile")
plt.ylabel("weight difference")
plt.show()

plt.plot(wgts_difference[0:10000:1000], '-o')
plt.title("every 10's Percentile ")
plt.grid()
plt.xlabel("percentile")
plt.ylabel("weight difference")
plt.show()
```

```


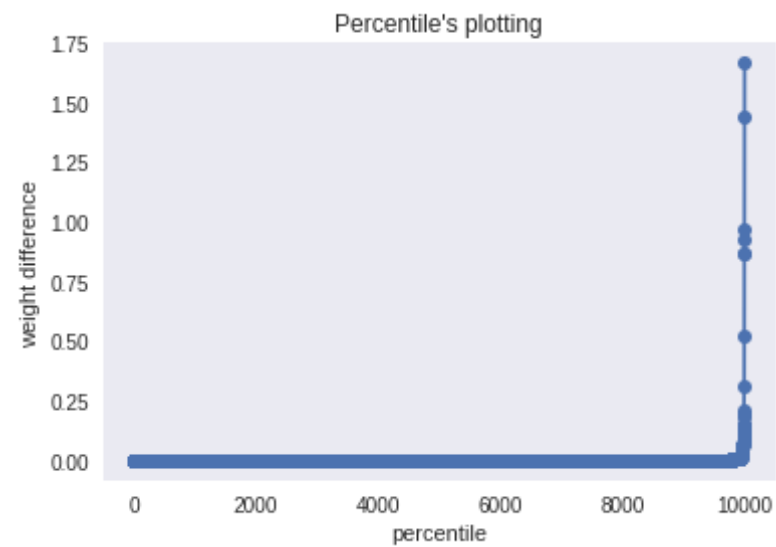
plt.plot(wgts_difference[0:1000:100], '-o')
plt.title("Percentile90")
plt.grid()
plt.xlabel("percentile")
plt.ylabel("weight difference")
plt.show()

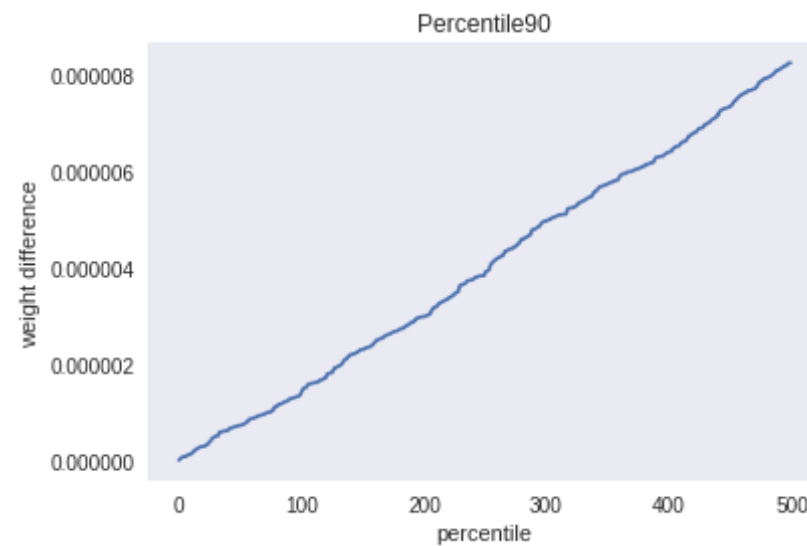

plt.plot(wgts_difference[0:500])
plt.title("Percentile90")
plt.grid()
plt.xlabel("percentile")
plt.ylabel("weight difference")
plt.show()

plt.plot(wgts_difference[0:10000], c='b')
plt.scatter(x=list(range(0,10000,2500)), y=wgts_difference[0:10000:2500], c='red', label="quantiles with 0.25 intervals")
# quantiles with 0.25 difference
plt.scatter(x=list(range(0,10000,1000)), y=wgts_difference[0:10000:1000], c='m', label = "quantiles with 0.05 intervals")
#for x,y in zip(list(range(0,100,25)), wgts_difference[0:100:25]):
#    plt.annotate(s="({} , {})".format(x,y), xy=(x,y))

plt.plot(wgts_difference[0:5000], c='b')
plt.scatter(x=list(range(0,5000,1250)), y=wgts_difference[0:5000:1250], c='red', label="quantiles with 0.25 intervals")
# quantiles with 0.25 difference
plt.scatter(x=list(range(0,5000,1250)), y=wgts_difference[0:5000:1250], c='m', label = "quantiles with 0.05 intervals")
#for x,y in zip(list(range(0,100,25)), wgts_difference[0:100:25]):
#    plt.annotate(s="({} , {})".format(x,y), xy=(x,y))

```





```
In [113]: wgts_difference.sort()
          wgts_difference[0:len(wgts_difference):1000]
```




```
Out[113]: array([2.95507520e-08, 1.61335340e-05, 3.20651557e-05, 4.95881036e-05,
                7.17487326e-05, 9.83537415e-05, 1.34006676e-04, 1.87983163e-04,
                2.90350526e-04, 6.02159641e-04])
```

```
In [114]: print(type(wgts_difference))
sns.boxplot(wgts_difference)
plt.show()

from numpy import percentile
val= percentile(wgts_difference,[25,50,75])
val

np.sort(wgts_difference).size

<class 'numpy.ndarray'>
```



```
Out[114]: 10000
```

[5.1.3] Feature Importance on BOW, SET 1

[5.1.3.1] Top 10 important features of positive class from SET 1

```
In [115]: feature_names=count_vect.get_feature_names()
          coefs=sorted(zip(clf1.coef_[0],feature_names))

          top20Negative=coefs[:20]
          top20Postive=coefs[::-1][:20]

          res_neg=pd.DataFrame(top20Negative,columns=['Features','Values'])
          res_pos=pd.DataFrame(top20Postive,columns=['Features','Values'])
          pd.concat([res_neg,res_pos],axis=1)
```

Out[115]:

	Features	Values	Features	Values
0	-0.261300	disappoint	0.497705	great
1	-0.182547	worst	0.425680	love
2	-0.163124	return	0.372925	best
3	-0.155267	terribl	0.295090	delici
4	-0.150823	thought	0.276514	good
5	-0.148989	aw	0.241148	excel
6	-0.148580	money	0.236529	perfect
7	-0.139838	horribl	0.220495	favorit
8	-0.136285	unfortun	0.214225	nice
9	-0.126837	threw	0.166982	wonder
10	-0.124238	stale	0.158020	amaz
11	-0.123137	bad	0.153355	easi
12	-0.122603	bland	0.151370	awesom
13	-0.121816	didnt	0.147459	use
14	-0.116762	wast	0.144988	glad
15	-0.107820	did	0.142276	addict
16	-0.106468	tast	0.139096	alway

	Features	Values	Features	Values
17	-0.099864	poor	0.137252	happi
18	-0.099568	hope	0.137182	tasti
19	-0.099322	refund	0.127082	thank

[5.1.3.2] Top 10 important features of negative class from SET 1

In [0]: *# Please write all the code with proper documentation*

[5.2] Logistic Regression on TFIDF, SET 2

[5.2.1] Applying Logistic Regression with L1 regularization on TFIDF, SET 2

```
In [0]: tf_idf_vect = TfidfVectorizer(ngram_range=(1,2), min_df=10, max_df=0.95
,stop_words='english',max_features=50000 )
tf_idf_vect.fit(X_train)
print("some sample features(unique words in the corpus)",tf_idf_vect.ge
t_feature_names()[0:10])
print('='*50)

X_train_tfidf= tf_idf_vect.transform(X_train)
print("the type of count vectorizer ",type(X_train_tfidf))
print("the shape of out text TFIDF vectorizer ",X_train_tfidf.get_shape
())
print("the number of unique words including both unigrams and bigrams "
, X_train_tfidf.get_shape()[1])

X_test_tfidf = tf_idf_vect.transform(X_test)
print("the type of count vectorizer ",type(X_test_tfidf))
print("the shape of out text TFIDF vectorizer ",X_test_tfidf.get_shape
())
print("the number of unique words ", X_test_tfidf.get_shape()[1])
```

```
some sample features(unique words in the corpus) ['aback', 'abandon',  
'abc', 'abdomin', 'abil', 'abl', 'abl amazon', 'abl ani', 'abl anywhe  
r', 'abl break']
```

```
=====
```

```
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>  
the shape of out text TFIDF vectorizer (75000, 38379)  
the number of unique words including both unigrams and bigrams 38379  
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>  
the shape of out text TFIDF vectorizer (25000, 38379)  
the number of unique words 38379
```

```
In [0]: from sklearn.preprocessing import StandardScaler  
scaler = StandardScaler(with_mean=False)  
scaler.fit(X_train_tfidf)  
X_train_tfidf=scaler.transform(X_train_tfidf)  
  
X_test_tfidf=scaler.transform(X_test_tfidf)
```

```
In [0]: from sklearn.linear_model import LogisticRegression  
  
clf=LogisticRegression(penalty='l1')  
param_grid={'C':[1000,100,10,1,0.1,0.001,0.0001]}  
#timeseriessplit=TimeSeriesSplit(n_splits=10)  
gcv=GridSearchCV(clf,param_grid,cv=5,scoring='roc_auc')  
gcv.fit(X_train_tfidf,y_train)  
print(gcv.best_params_)  
print(gcv.best_score_)  
  
{'C': 0.1}  
0.9292712996949454
```


```
In [0]: hyper_parameters=gcv.get_params()['param_grid']['C']  
train_scores=gcv.cv_results_['mean_train_score'].tolist()  
test_scores=gcv.cv_results_['mean_test_score'].tolist()  
  
print(hyper_parameters)  
print(test_scores)  
print(train_scores)
```

```
plt.plot( hyper_parameters ,train_scores , label='Train plot')
plt.plot( hyper_parameters ,test_scores , label='Test plot')
plt.xlabel("hyper_parameters (C)")
plt.ylabel("Model performance")
```

```
plt.legend()
```

```
[1000, 100, 10, 1, 0.1, 0.001, 0.0001]
[0.8909053304983988, 0.9016360483725683, 0.9178976054051204, 0.92221051
00849963, 0.9292712996949454, 0.8672104253999816, 0.5]
[1.0, 1.0, 1.0, 1.0, 0.999998432394815, 0.8695069653061985, 0.5]
```

Out[0]: <matplotlib.legend.Legend at 0x7fe9a597dc18>



```
In [0]: from sklearn.metrics import roc_auc_score
from sklearn.metrics import auc
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
```

```

from sklearn.metrics import f1_score

clf1=LogisticRegression(C=0.1,penalty='l1')
clf1.fit(X_train_tfidf,y_train)
pred_train=clf1.predict(X_train_tfidf)
pred=clf1.predict(X_test_tfidf)

print("Accuracy Score : ",accuracy_score(y_test,pred)*100)
print("Precision Score : ",precision_score(y_test,pred)*100)
print("Recall Score : ",recall_score(y_test,pred)*100)
print("F1 Score : ",f1_score(y_test,pred)*100)

print(" ")
print("Classification Report")
print(classification_report(y_test,pred))
print(" ")

fpr_train,tpr_train,thresholds_train=roc_curve(y_train,pred_train)
print("AUC Score for train data :",metrics.auc(fpr_train,tpr_train))

fpr,tpr,thresholds=roc_curve(y_test,pred)
print("AUC Score for test data :",metrics.auc(fpr,tpr))

print(" ")

#y_true = # ground truth labels
#y_probas = # predicted probabilities generated by sklearn classifier
#skplt.metrics.plot_roc_curve(y_true, y_probas)
#plt.show()

plt.figure()
lw = 2
plt.plot(fpr, tpr, color='red',
         lw=lw,label='test')
plt.plot(fpr_train, tpr_train, color='darkorange',

```

```

        lw=lw,label='train')
plt.plot([0, 1], [0, 1], color='navy', lw=lw, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic example')
plt.legend(loc="lower right")
plt.show()

print("      ")

tn, fp, fn, tp=confusion_matrix(y_test,pred).ravel()
print("""
TrueNegative : {}
FalsePositive : {}
FalseNegative : {}
TruePositive : {}""".format(tn, fp, fn, tp))
print("      ")
print("      ")

confusionmatrix_DF=pd.DataFrame(confusion_matrix(y_test,pred),columns=[
'0','1'],index=['0','1'])
sns.heatmap(confusionmatrix_DF,annot=True,fmt='g',cmap='viridis')
plt.title("Confusion matrix ")
plt.show()

```


Accuracy Score : 91.444
Precision Score : 94.64103278536268
Recall Score : 95.38965597940557
F1 Score : 95.01386978717453

Classification Report				
	precision	recall	f1-score	support
0	0.72	0.68	0.70	3635
1	0.95	0.95	0.95	21365
micro avg	0.91	0.91	0.91	25000


macro avg	0.83	0.82	0.82	25000
weighted avg	0.91	0.91	0.91	25000

AUC Score for train data : 0.9995970630372493

AUC Score for test data : 0.8182137544499852



TrueNegative : 2481
FalsePositive : 1154
FalseNegative : 985
TruePositive : 20380



[5.2.2] Applying Logistic Regression with L2 regularization on TFIDF, SET 2

```
In [0]: from sklearn.linear_model import LogisticRegression
```

```
clf=LogisticRegression(penalty='l2')
param_grid={'C':[1000,100,10,1,0.1,0.001,0.0001]}
#timeseriessplit=TimeSeriesSplit(n_splits=10)
gcv=GridSearchCV(clf,param_grid,cv=5,scoring='roc_auc')
gcv.fit(X_train_tfidf,y_train)
print(gcv.best_params_)
print(gcv.best_score_)
```

```
{'C': 0.0001}
0.9386103426546142
```

```
In [0]: hyper_parameters=gcv.get_params()['param_grid']['C']
train_scores=gcv.cv_results_['mean_train_score'].tolist()
test_scores=gcv.cv_results_['mean_test_score'].tolist()

print(hyper_parameters)
```

```

print(test_scores)
print(train_scores)

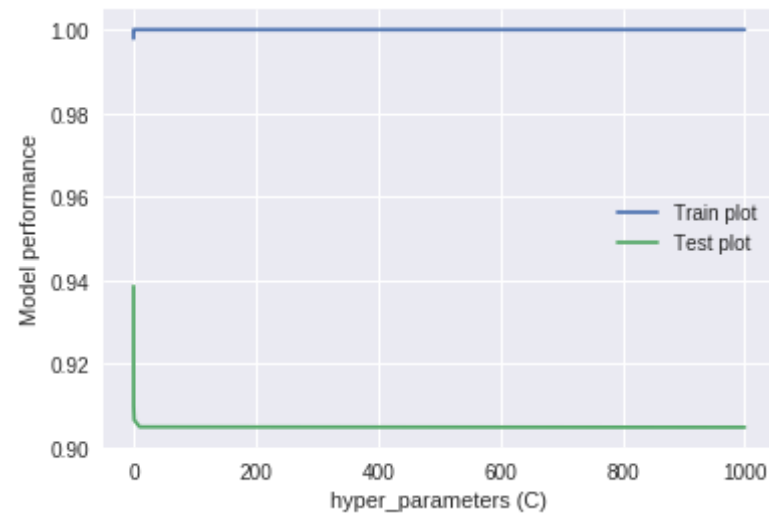
plt.plot( hyper_parameters ,train_scores , label='Train plot')
plt.plot( hyper_parameters ,test_scores , label='Test plot')
plt.xlabel("hyper_parameters (C)")
plt.ylabel("Model performance")

plt.legend()

[1000, 100, 10, 1, 0.1, 0.001, 0.0001]
[0.9048738818353523, 0.9049524106876847, 0.904967451186365, 0.906631545
2519695, 0.9098220252645793, 0.926345752661734, 0.9386103426546142]
[1.0, 1.0, 1.0, 1.0, 1.0, 0.9999967420917926, 0.9978873148839904]

```

Out[0]: <matplotlib.legend.Legend at 0x7fe9af91ae48>



```

In [0]: from sklearn.metrics import roc_auc_score
from sklearn.metrics import auc
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report

```

```

from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
from sklearn.metrics import f1_score

clf1=LogisticRegression(C=0.0001,penalty='l2')
clf1.fit(X_train_tfidf,y_train)
pred_train=clf1.predict(X_train_tfidf)
pred=clf1.predict(X_test_tfidf)

print("Accuracy Score : ",accuracy_score(y_test,pred)*100)
print("Precision Score : ",precision_score(y_test,pred)*100)
print("Recall Score : ",recall_score(y_test,pred)*100)
print("F1 Score : ",f1_score(y_test,pred)*100)

print(" ")
print("Classification Report")
print(classification_report(y_test,pred))
print(" ")

fpr_train,tpr_train,thresholds_train=roc_curve(y_train,pred_train)
print("AUC Score for train data :",metrics.auc(fpr_train,tpr_train))

fpr,tpr,thresholds=roc_curve(y_test,pred)
print("AUC Score for test data :",metrics.auc(fpr,tpr))

print(" ")

#y_true = # ground truth labels
#y_probas = # predicted probabilities generated by sklearn classifier
#skplt.metrics.plot_roc_curve(y_true, y_probas)
#plt.show()

plt.figure()
lw = 2
plt.plot(fpr, tpr, color='red',

```

```

        lw=lw,label='test')
plt.plot(fpr_train, tpr_train, color='darkorange',
        lw=lw,label='train')
plt.plot([0, 1], [0, 1], color='navy', lw=lw, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic example')
plt.legend(loc="lower right")
plt.show()

print(" ")

tn, fp, fn, tp=confusion_matrix(y_test,pred).ravel()
print("""
TrueNegative : {}
FalsePositive : {}
FalseNegative : {}
TruePositive : {}""".format(tn, fp, fn, tp))
print(" ")
print(" ")

confusionmatrix_DF=pd.DataFrame(confusion_matrix(y_test,pred),columns=[
'0','1'],index=['0','1'])
sns.heatmap(confusionmatrix_DF,annot=True,fmt='g',cmap='viridis')
plt.title("Confusion matrix ")
plt.show()

```


Accuracy Score : 91.776
Precision Score : 92.6831424908263
Recall Score : 98.12309852562603
F1 Score : 95.32557293561294

Classification Report				
	precision	recall	f1-score	support
0	0.83	0.54	0.66	3635
1	0.93	0.98	0.95	21365


micro avg	0.92	0.92	0.92	25000
macro avg	0.88	0.76	0.81	25000
weighted avg	0.91	0.92	0.91	25000

AUC Score for train data : 0.9125294110148021

AUC Score for test data : 0.7629676246776487



TrueNegative : 1980
 FalsePostive : 1655
 FalseNegative : 401
 TruePostive : 20964



[5.2.3] Feature Importance on TFIDF, SET 2

[5.2.3.1] Top 10 important features of positive class from SET 2

```
In [0]: feature_names=tf_idf_vect.get_feature_names()
        coefs=sorted(zip(clf1.coef_[0],feature_names))

        top20Negative=coefs[:20]
        top20Postive=coefs[::-1][:20]

        res_neg=pd.DataFrame(top20Negative,columns=['Features','Values'])
        res_pos=pd.DataFrame(top20Postive,columns=['Features','Values'])
        pd.concat([res_neg,res_pos],axis=1)
```

Out[0]:

	Features	Values	Features	Values
0	-0.071519	disappoint	0.122651	great
1	-0.067153	veri disappoint	0.121623	love

	Features	Values	Features	Values
2	-0.060787	worst	0.094357	best
3	-0.051640	terribl	0.080949	good
4	-0.050525	aw	0.075217	delici
5	-0.047269	return	0.059922	perfect
6	-0.045456	horribl	0.059878	excel
7	-0.042103	threw	0.058697	favorit
8	-0.041307	wors	0.052519	nice
9	-0.041176	wont buy	0.052321	use
10	-0.040446	wast money	0.051268	wonder
11	-0.040267	bland	0.048090	easi
12	-0.040092	wast	0.046017	high recommend
13	-0.039738	disgust	0.045918	make
14	-0.037786	money	0.042632	enjoy
15	-0.037536	tast like	0.041551	thank
16	-0.037288	stale	0.041076	alway
17	-0.037015	unfortun	0.040092	tasti
18	-0.036412	tasteless	0.039935	amaz
19	-0.033395	refund	0.039808	awesom

[5.2.3.2] Top 10 important features of negative class from SET 2

In [0]: `# Please write all the code with proper documentation`

[5.3] Logistic Regression on AVG W2V, SET 3

[5.3.1] Applying Logistic Regression with L1 regularization on AVG W2V SET 3

```
In [0]: # Train your own Word2Vec model using your own text corpus
i=0
list_of_sentence=[]
for sentence in X_train:
    list_of_sentence.append(sentence.split())

#####

is_your_ram_gt_16g=False
want_to_use_google_w2v = False
want_to_train_w2v = True

if want_to_train_w2v:
    # min_count = 5 considers only words that occurred at least 5 times
    w2v_model=Word2Vec(list_of_sentence,min_count=5,size=50, workers=4)
    print(w2v_model.wv.most_similar('great'))
    print('='*50)
    print(w2v_model.wv.most_similar('worst'))

elif want_to_use_google_w2v and is_your_ram_gt_16g:
    if os.path.isfile('GoogleNews-vectors-negative300.bin'):
        w2v_model=KeyedVectors.load_word2vec_format('GoogleNews-vectors-
-negative300.bin', binary=True)
        print(w2v_model.wv.most_similar('great'))
        print(w2v_model.wv.most_similar('worst'))
    else:
        print("you don't have google's word2vec file, keep want_to_train_w2v = True, to train your own w2v ")

#####

w2v_words = list(w2v_model.wv.vocab)
print("number of words that occurred minimum 5 times ",len(w2v_words))
```



```

print("sample words ", w2v_words[0:50])

#*****
#*****

# average Word2Vec
# compute average word2vec for each review.
X_train_AvgW2V_100000 = []; # the avg-w2v for each sentence/review is s
tored in this list
for sent in tqdm(list_of_sentence): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length 50, yo
u might need to change this to 300 if you use google's w2v
    cnt_words = 0; # num of words with a valid vector in the sentence/re
view
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
        if cnt_words != 0:
            sent_vec /= cnt_words
    X_train_AvgW2V_100000.append(sent_vec)
print(len(X_train_AvgW2V_100000))
print(len(X_train_AvgW2V_100000[0]))

```

```

0%|          | 92/75000 [00:00<01:21, 915.01it/s]

```

```

[('terrif', 0.8637382388114929), ('awesom', 0.8417178392410278), ('exce
l', 0.8403081297874451), ('wonder', 0.8175906538963318), ('fantast', 0.
8126950263977051), ('good', 0.8095090389251709), ('perfect', 0.80067425
96626282), ('fabul', 0.7345156669616699), ('decent', 0.70958995819091
8), ('nice', 0.6975454688072205)]

```

```

=====

```

```

[('greatest', 0.7696366310119629), ('best', 0.7653523087501526), ('tast
iest', 0.6937639713287354), ('nicest', 0.6832467913627625), ('closest',
0.6737667918205261), ('disgust', 0.6405009627342224), ('nastiest', 0.63
55428695678711), ('hottest', 0.5996732115745544), ('finest', 0.58700078
72581482), ('horribl', 0.5557085275650024)]

```

```

number of words that occured minimum 5 times 11759

```

```
sample words ['order', 'this', 'the', 'tastiest', 'oatmeal', 'have',  
'ever', 'eaten', 'far', 'superior', 'organ', 'been', 'purchas', 'natu  
r', 'section', 'local', 'groceri', 'store', 'better', 'flavor', 'and',  
'textur', 'also', 'servic', 'was', 'great', 'summertim', 'has', 'gone',  
'but', 'product', 'still', 'tast', 'hot', 'drink', 'for', 'summer', 'we  
ll', 'not', 'know', 'could', 'surviv', 'heat', 'without', 'cooler', 'ic  
i', 'slush', 'unfortun', 'all', 'friend']
```

```
100%|██████████| 75000/75000 [02:07<00:00, 589.43it/s]
```

```
75000
```

```
50
```

```
In [0]: import pickle  
with open('X_train_AvgW2V.pkl', 'wb') as f:  
    pickle.dump(X_train_AvgW2V, f)
```

```
In [0]: import pickle  
with open('X_train_AvgW2V_100000.pkl', 'wb') as f:  
    pickle.dump(X_train_AvgW2V_100000, f)
```

```
In [0]: #*****  
*****  
  
i=0  
list_of_sentence_test=[]  
for sentence in X_test:  
    list_of_sentence_test.append(sentence.split())  
  
#*****  
*****  
  
# average Word2Vec  
# compute average word2vec for each review.  
X_test_AvgW2V_100000 = []; # the avg-w2v for each sentence/review is st  
ored in this list  
for sent in tqdm(list_of_sentence_test): # for each review/sentence
```

```

sent_vec = np.zeros(50) # as word vectors are of zero length 50, yo
u might need to change this to 300 if you use google's w2v
cnt_words = 0; # num of words with a valid vector in the sentence/re
view
for word in sent: # for each word in a review/sentence
    if word in w2v_words:
        vec = w2v_model.wv[word]
        sent_vec += vec
        cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
X_test_AvgW2V_100000.append(sent_vec)
print(len(X_test_AvgW2V_100000))
print(len(X_test_AvgW2V_100000[0]))

```

```

100%|██████████| 25000/25000 [00:43<00:00, 579.18it/s]

```

```

25000
50

```

```

In [0]: from google.colab import files
files.download('X_train_AvgW2V_100000.pkl')

```

```

In [0]: import pickle
with open('X_test_AvgW2V_100000.pkl', 'wb') as f:
    pickle.dump(X_test_AvgW2V_100000, f)

```

```

In [0]: from google.colab import files
files.download('X_test_AvgW2V_100000.pkl')

```

```

In [0]: from sklearn.preprocessing import StandardScaler
scaler = StandardScaler(with_mean=False)
scaler.fit(X_train_AvgW2V_100000)
X_train_AvgW2V=scaler.transform(X_train_AvgW2V_100000)

X_test_AvgW2V=scaler.transform(X_test_AvgW2V_100000)

```

```
In [0]: from sklearn.linear_model import LogisticRegression
```

```
clf=LogisticRegression(penalty='l1')
param_grid={'C':[1000,100,10,5,1,0.5,0.1,0.001,0.0001]}
#timeseriessplit=TimeSeriesSplit(n_splits=10)
gcv=GridSearchCV(clf,param_grid,cv=5,scoring='roc_auc')
gcv.fit(X_train_AvgW2V,y_train)
print(gcv.best_params_)
print(gcv.best_score_)
```

```
{'C': 0.1}
0.9007299333163759
```

```
In [0]: hyper_parameters=gcv.get_params()['param_grid']['C']
train_scores=gcv.cv_results_['mean_train_score'].tolist()
test_scores=gcv.cv_results_['mean_test_score'].tolist()
```


```
print(hyper_parameters)
print(test_scores)
print(train_scores)
```

```
plt.plot( hyper_parameters ,train_scores , label='Train plot')
plt.plot( hyper_parameters ,test_scores , label='Test plot')
plt.xlabel("hyper_parameters (C)")
plt.ylabel("Model performance")
```

```
plt.legend()
```

```
[1000, 100, 10, 5, 1, 0.5, 0.1, 0.001, 0.0001]
[0.9007167027794262, 0.9007180144130086, 0.9007181265108707, 0.90071928
35458553, 0.9007202587848292, 0.9007170469938147, 0.9007299333163759,
0.8782740824196674, 0.6711832154829824]
[0.9014580585714148, 0.9014582111242412, 0.9014587867253511, 0.90145940
21766745, 0.9014579165038056, 0.9014517503699985, 0.901458403119588, 0.
8788911387273874, 0.6712408567467387]
```

```
Out[0]: <matplotlib.legend.Legend at 0x7fe9af3d6eb8>
```



```
In [0]: from sklearn.metrics import roc_auc_score
from sklearn.metrics import auc
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
from sklearn.metrics import f1_score

clf1=LogisticRegression(C=0.1,penalty='l1')
clf1.fit(X_train_AvgW2V,y_train)
pred_train=clf1.predict(X_train_AvgW2V)
pred=clf1.predict(X_test_AvgW2V)

print("Accuracy Score : ",accuracy_score(y_test,pred)*100)
print("Precision Score : ",precision_score(y_test,pred)*100)
print("Recall Score : ",recall_score(y_test,pred)*100)
print("F1 Score : ",f1_score(y_test,pred)*100)

print(" ")
print("Classification Report")
```

```

print(classification_report(y_test,pred))
print(" ")

fpr_train,tpr_train,thresholds_train=roc_curve(y_train,pred_train)
print("AUC Score for train data :",metrics.auc(fpr_train,tpr_train))

fpr,tpr,thresholds=roc_curve(y_test,pred)
print("AUC Score for test data :",metrics.auc(fpr,tpr))

print(" ")

#y_true = # ground truth labels
#y_probas = # predicted probabilities generated by sklearn classifier
#skplt.metrics.plot_roc_curve(y_true, y_probas)
#plt.show()

plt.figure()
lw = 2
plt.plot(fpr, tpr, color='red',
         lw=lw,label='test')
plt.plot(fpr_train, tpr_train, color='darkorange',
         lw=lw,label='train')
plt.plot([0, 1], [0, 1], color='navy', lw=lw, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic example')
plt.legend(loc="lower right")
plt.show()

print(" ")

tn, fp, fn, tp=confusion_matrix(y_test,pred).ravel()
print("""
TrueNegative : {}

```

```

FalsePositive : {}
FalseNegative : {}
TruePositive : {}"".format(tn, fp, fn, tp))
print(" ")
print(" ")


confusionmatrix_DF=pd.DataFrame(confusion_matrix(y_test,pred),columns=[
'0','1'],index=['0','1'])
sns.heatmap(confusionmatrix_DF,annot=True,fmt='g',cmap='viridis')
plt.title("Confusion matrix ")
plt.show()

```


Accuracy Score : 88.74
 Precision Score : 90.73699929725932
 Recall Score : 96.69553007254856
 F1 Score : 93.6215530328779

Classification Report					
	precision	recall	f1-score	support	
0	0.68	0.42	0.52	3635	
1	0.91	0.97	0.94	21365	
micro avg	0.89	0.89	0.89	25000	
macro avg	0.80	0.69	0.73	25000	
weighted avg	0.87	0.89	0.88	25000	

AUC Score for train data : 0.6994174199573306
 AUC Score for test data : 0.6933813642554526



TrueNegative : 1526
FalsePositive : 2109
FalseNegative : 706
TruePositive : 20659



[5.3.2] Applying Logistic Regression with L2 regularization on AVG W2V, SET 3

In [0]: `from sklearn.linear_model import LogisticRegression`

```
clf=LogisticRegression(penalty='l2')
param_grid={'C':[1000,100,10,5,1,0.5,0.1,0.001,0.0001]}
#timeseriessplit=TimeSeriesSplit(n_splits=10)
gcv=GridSearchCV(clf,param_grid,cv=5,scoring='roc_auc')
gcv.fit(X_train_AvgW2V,y_train)
print(gcv.best_params_)
print(gcv.best_score_)
```

```
{'C': 0.1}
0.9007230793376875
```

In [0]: `hyper_parameters=gcv.get_params()['param_grid']['C']`
`train_scores=gcv.cv_results_['mean_train_score'].tolist()`
`test_scores=gcv.cv_results_['mean_test_score'].tolist()`

```
print(hyper_parameters)
print(test_scores)
print(train_scores)
```


```
plt.plot( hyper_parameters ,train_scores , label='Train plot')
plt.plot( hyper_parameters ,test_scores , label='Test plot')
plt.xlabel("hyper_parameters (C)")
plt.ylabel("Model performance")
```

```
plt.legend()
```

```
[1000, 100, 10, 5, 1, 0.5, 0.1, 0.001, 0.0001]
[0.9007046960046837, 0.9007049624469834, 0.9007073262931394, 0.90070835
69485211, 0.9007160648083817, 0.9007200974046736, 0.9007230793376875,
0.9007230793376875]
```

```
0.8988671134516509, 0.882827981371257]
[0.9014466984609755, 0.901446754567538, 0.9014495941602446, 0.901449513
5941442, 0.9014572699788624, 0.9014608961262898, 0.9014580909601155, 0.
8994398345547424, 0.8832301217866835]
```

Out[0]: <matplotlib.legend.Legend at 0x7fe9aecddac8>



```
In [0]: from sklearn.metrics import roc_auc_score
from sklearn.metrics import auc
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
from sklearn.metrics import f1_score
```

```
clf1=LogisticRegression(C=0.1,penalty='l2')
clf1.fit(X_train_AvgW2V,y_train)
pred_train=clf1.predict(X_train_AvgW2V)
pred=clf1.predict(X_test_AvgW2V)
```

```

print("Accuracy Score : ",accuracy_score(y_test,pred)*100)
print("Precision Score : ",precision_score(y_test,pred)*100)
print("Recall Score : ",recall_score(y_test,pred)*100)
print("F1 Score : ",f1_score(y_test,pred)*100)

print(" ")
print("Classification Report")
print(classification_report(y_test,pred))
print(" ")

fpr_train,tpr_train,thresholds_train=roc_curve(y_train,pred_train)
print("AUC Score for train data :",metrics.auc(fpr_train,tpr_train))

fpr,tpr,thresholds=roc_curve(y_test,pred)
print("AUC Score for test data :",metrics.auc(fpr,tpr))

print(" ")

#y_true = # ground truth labels
#y_probas = # predicted probabilities generated by sklearn classifier
#skplt.metrics.plot_roc_curve(y_true, y_probas)
#plt.show()

plt.figure()
lw = 2
plt.plot(fpr, tpr, color='red',
         lw=lw,label='test')
plt.plot(fpr_train, tpr_train, color='darkorange',
         lw=lw,label='train')
plt.plot([0, 1], [0, 1], color='navy', lw=lw, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic example')
plt.legend(loc="lower right")
plt.show()

```

```

print("      ")

tn, fp, fn, tp=confusion_matrix(y_test,pred).ravel()
print("""
TrueNegative : {}
FalsePositive : {}
FalseNegative : {}
TruePositive : {}""".format(tn, fp, fn, tp))
print("      ")
print("      ")


confusionmatrix_DF=pd.DataFrame(confusion_matrix(y_test,pred),columns=[
'0','1'],index=['0','1'])
sns.heatmap(confusionmatrix_DF,annot=True,fmt='g',cmap='viridis')
plt.title("Confusion matrix ")
plt.show()

```


Accuracy Score : 88.764
Precision Score : 90.77883263009845
Recall Score : 96.6721273110227
F1 Score : 93.63283994831923

Classification Report				
	precision	recall	f1-score	support
0	0.68	0.42	0.52	3635
1	0.91	0.97	0.94	21365
micro avg	0.89	0.89	0.89	25000
macro avg	0.80	0.69	0.73	25000
weighted avg	0.88	0.89	0.88	25000

AUC Score for train data : 0.7005523554218793
AUC Score for test data : 0.6947774178480983



TrueNegative : 1537
FalsePositive : 2098
FalseNegative : 711
TruePositive : 20654



[5.4] Logistic Regression on TFIDF W2V, SET 4

[5.4.1] Applying Logistic Regression with L1 regularization on TFIDF W2V, SET 4

```
In [0]: # S = ["abc def pqr", "def def def abc", "pqr pqr def"]
model = TfidfVectorizer(ngram_range=(1,2), min_df=10, max_df=0.95, stop_
words='english', max_features=5000 )
tf_idf_matrix = model.fit_transform(X_train)
# we are converting a dictionary with word as a key, and the idf as a v
alue
dictionary = dict(zip(model.get_feature_names(), list(model.idf_)))

#*****
#*****

# TF-IDF weighted Word2Vec
tfidf_feat = model.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and ce
ll_val = tfidf

X_train_Avgtfidf_100000 = []; # the tfidf-w2v for each sentence/review
is stored in this list
row=0;
for sent in list_of_sentence: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum = 0; # num of words with a valid vector in the sentence/r
review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
#             tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
# to reduce the computation we are
# dictionary[word] = idf value of word in whole corpus
```

```

        # sent.count(word) = tf value of word in this review
        tf_idf = dictionary[word]*(sent.count(word)/len(sent))
        sent_vec += (vec * tf_idf)
        weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    X_train_Avgtfidf_100000.append(sent_vec)
    row += 1

```

```

In [0]: *****
*****

X_test_Avgtfidf_100000 = []; # the tfidf-w2v for each sentence/review is stored in this list
row=0;
for sent in list_of_sentence_test: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
            #
            # tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
            # to reduce the computation we are
            # dictionary[word] = idf value of word in whole corpus
            # sent.count(word) = tf value of word in this review
            tf_idf = dictionary[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    X_test_Avgtfidf_100000.append(sent_vec)
    row += 1

```

```

3%|| | 831/25000 [00:37<03:33, 113.03it/s]

```

```

In [0]: import pickle
with open('X_train_Avgtfidf_100000.pkl', 'wb') as f:

```

```
pickle.dump(X_train_Avgtfidf_100000, f)
```

```
In [0]: import pickle  
with open('X_test_Avgtfidf_100000.pkl', 'wb') as f:  
    pickle.dump(X_test_Avgtfidf_100000, f)
```

```
In [0]: from google.colab import files  
files.download('X_train_Avgtfidf_100000.pkl')
```

```
In [0]: from google.colab import files  
files.download('X_test_Avgtfidf_100000.pkl')
```

```
In [0]: from sklearn.preprocessing import StandardScaler  
scaler = StandardScaler(with_mean=False)  
scaler.fit(X_train_Avgtfidf_100000)  
X_train_Avgtfidf=scaler.transform(X_train_Avgtfidf_100000)  
  
X_test_Avgtfidf=scaler.transform(X_test_Avgtfidf_100000)
```

```
In [0]: print(X_train_Avgtfidf.shape, len(y_train), X_test_Avgtfidf.shape, len(y_t  
est))  
  
(75000, 50) 75000 (25000, 50) 25000
```

```
In [0]: from sklearn.linear_model import LogisticRegression  
  
clf=LogisticRegression(penalty='l1')  
param_grid={'C':[1000,100,10,5,1,0.5,0.1,0.001,0.0001]}  
#timeseriessplit=TimeSeriesSplit(n_splits=10)  
gcv=GridSearchCV(clf,param_grid,cv=5,scoring='roc_auc')  
gcv.fit(X_train_Avgtfidf,y_train)  
print(gcv.best_params_)  
print(gcv.best_score_)  
  
{'C': 0.5}  
0.8634936283951694
```



```

In [0]: hyper_parameters=gcv.get_params()['param_grid']['C']
train_scores=gcv.cv_results_['mean_train_score'].tolist()
test_scores=gcv.cv_results_['mean_test_score'].tolist()

print(hyper_parameters)
print(test_scores)
print(train_scores)


plt.plot( hyper_parameters ,train_scores , label='Train plot')
plt.plot( hyper_parameters ,test_scores , label='Test plot')
plt.xlabel("hyper_parameters (C)")
plt.ylabel("Model performance")

plt.legend()

[1000, 100, 10, 5, 1, 0.5, 0.1, 0.001, 0.0001]
[0.8634130341529095, 0.8634134832056587, 0.8634180844298626, 0.86342183
02513165, 0.8634629300248184, 0.8634936283951694, 0.8634162350503543,
0.8152802608557099, 0.6386246100716295]
[0.8646048120674523, 0.8646051951826704, 0.8646106962769953, 0.86461335
87574667, 0.8646545997067424, 0.8646792115276515, 0.8645850224511005,
0.8160322620536167, 0.6386806669634713]

```

Out[0]: <matplotlib.legend.Legend at 0x7fe9af210c88>



```
In [0]: from sklearn.metrics import roc_auc_score
from sklearn.metrics import auc
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
from sklearn.metrics import f1_score

clf1=LogisticRegression(C=0.5,penalty='l1')
clf1.fit(X_train_Avgtfidf,y_train)
pred_train=clf1.predict(X_train_Avgtfidf)
pred=clf1.predict(X_test_Avgtfidf)

print("Accuracy Score : ",accuracy_score(y_test,pred)*100)
print("Precision Score : ",precision_score(y_test,pred)*100)
print("Recall Score : ",recall_score(y_test,pred)*100)
print("F1 Score : ",f1_score(y_test,pred)*100)

print(" ")
print("Classification Report")
```

```

print(classification_report(y_test,pred))
print(" ")

fpr_train,tpr_train,thresholds_train=roc_curve(y_train,pred_train)
print("AUC Score for train data :",metrics.auc(fpr_train,tpr_train))

fpr,tpr,thresholds=roc_curve(y_test,pred)
print("AUC Score for test data :",metrics.auc(fpr,tpr))

print(" ")

#y_true = # ground truth labels
#y_probas = # predicted probabilities generated by sklearn classifier
#skplt.metrics.plot_roc_curve(y_true, y_probas)
#plt.show()

plt.figure()
lw = 2
plt.plot(fpr, tpr, color='red',
         lw=lw,label='test')
plt.plot(fpr_train, tpr_train, color='darkorange',
         lw=lw,label='train')
plt.plot([0, 1], [0, 1], color='navy', lw=lw, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic example')
plt.legend(loc="lower right")
plt.show()

print(" ")

tn, fp, fn, tp=confusion_matrix(y_test,pred).ravel()
print("""
TrueNegative : {}

```

```

FalsePositive : {}
FalseNegative : {}
TruePositive : {}"".format(tn, fp, fn, tp))
print(" ")
print(" ")


confusionmatrix_DF=pd.DataFrame(confusion_matrix(y_test,pred),columns=[
'0','1'],index=['0','1'])
sns.heatmap(confusionmatrix_DF,annot=True,fmt='g',cmap='viridis')
plt.title("Confusion matrix ")
plt.show()

```


Accuracy Score : 87.52799999999999
 Precision Score : 89.2898669307954
 Recall Score : 97.04657149543647
 F1 Score : 93.00677342663617

Classification Report					
	precision	recall	f1-score	support	
0	0.65	0.32	0.42	3635	
1	0.89	0.97	0.93	21365	
micro avg	0.88	0.88	0.88	25000	
macro avg	0.77	0.64	0.68	25000	
weighted avg	0.86	0.88	0.86	25000	

AUC Score for train data : 0.6485271285189526
 AUC Score for test data : 0.6431420734331659



TrueNegative : 1148
FalsePostive : 2487
FalseNegative : 631
TruePostive : 20734



[5.4.2] Applying Logistic Regression with L2 regularization on TFIDF W2V, SET 4

```
In [0]: from sklearn.linear_model import LogisticRegression
```

```
clf=LogisticRegression(penalty='l2')
param_grid={'C':[1000,100,10,5,1,0.5,0.1,0.001,0.0001]}
#timeseriessplit=TimeSeriesSplit(n_splits=10)
gcv=GridSearchCV(clf,param_grid,cv=5,scoring='roc_auc')
gcv.fit(X_train_Avgtfidf,y_train)
print(gcv.best_params_)
print(gcv.best_score_)
```

```
{'C': 0.1}
0.8635213967923081
```

```
In [0]: hyper_parameters=gcv.get_params()['param_grid']['C']
train_scores=gcv.cv_results_['mean_train_score'].tolist()
test_scores=gcv.cv_results_['mean_test_score'].tolist()
```

```

print(hyper_parameters)
print(test_scores)
print(train_scores)


plt.plot( hyper_parameters ,train_scores , label='Train plot')
plt.plot( hyper_parameters ,test_scores , label='Test plot')
plt.xlabel("hyper_parameters (C)")
plt.ylabel("Model performance")

plt.legend()

[1000, 100, 10, 5, 1, 0.5, 0.1, 0.001, 0.0001]
[0.8633956265099353, 0.8633952826912352, 0.8634009355048974, 0.86340985
6826655, 0.8634449191924536, 0.8634786199808072, 0.8635213967923081, 0.
8592823633143652, 0.8240068507338287]
[0.8645871843553113, 0.8645869507616778, 0.8645923646056921, 0.86459955
91362844, 0.8646347196807469, 0.8646650113597101, 0.8646980758164574,
0.8602040061553147, 0.8245984473005616]

```

Out[0]: <matplotlib.legend.Legend at 0x7fe9a8f2c4a8>



In [0]: `from sklearn.metrics import roc_auc_score`

```

from sklearn.metrics import auc
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
from sklearn.metrics import f1_score

clf1=LogisticRegression(C=0.1,penalty='l2')
clf1.fit(X_train_Avgtfidf,y_train)
pred_train=clf1.predict(X_train_Avgtfidf)
pred=clf1.predict(X_test_Avgtfidf)

print("Accuracy Score : ",accuracy_score(y_test,pred)*100)
print("Precision Score : ",precision_score(y_test,pred)*100)
print("Recall Score : ",recall_score(y_test,pred)*100)
print("F1 Score : ",f1_score(y_test,pred)*100)

print(" ")
print("Classification Report")
print(classification_report(y_test,pred))
print(" ")

fpr_train,tpr_train,thresholds_train=roc_curve(y_train,pred_train)
print("AUC Score for train data :",metrics.auc(fpr_train,tpr_train))

fpr,tpr,thresholds=roc_curve(y_test,pred)
print("AUC Score for test data :",metrics.auc(fpr,tpr))

print(" ")

#y_true = # ground truth labels
#y_probas = # predicted probabilities generated by sklearn classifier
#skplt.metrics.plot_roc_curve(y_true, y_probas)
#plt.show()

```



```

plt.figure()
lw = 2
plt.plot(fpr, tpr, color='red',
         lw=lw, label='test')
plt.plot(fpr_train, tpr_train, color='darkorange',
         lw=lw, label='train')
plt.plot([0, 1], [0, 1], color='navy', lw=lw, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic example')
plt.legend(loc="lower right")
plt.show()

print(" ")

tn, fp, fn, tp=confusion_matrix(y_test,pred).ravel()
print("""
TrueNegative : {}
FalsePostive : {}
FalseNegative : {}
TruePostive : {}""".format(tn, fp, fn, tp))
print(" ")
print(" ")

confusionmatrix_DF=pd.DataFrame(confusion_matrix(y_test,pred),columns=[
'0','1'],index=['0','1'])
sns.heatmap(confusionmatrix_DF,annot=True,fmt='g',cmap='viridis')
plt.title("Confusion matrix ")
plt.show()

```

```

Accuracy Score : 87.53999999999999
Precision Score : 89.2946343984153
Recall Score : 97.0559326000468
F1 Score : 93.0136586897526


```

Classification Report


	precision	recall	f1-score	support
0	0.65	0.32	0.42	3635
1	0.89	0.97	0.93	21365
micro avg	0.88	0.88	0.88	25000
macro avg	0.77	0.64	0.68	25000
weighted avg	0.86	0.88	0.86	25000

AUC Score for train data : 0.6477156491287481

AUC Score for test data : 0.6433264305380608



TrueNegative : 1149
 FalsePositive : 2486
 FalseNegative : 629
 TruePositive : 20736



[6] Conclusions

```
In [0]: from prettytable import PrettyTable

x = PrettyTable()

x.field_names = ["LogisticRegression with Different Vectorization", "penalty" , "C-value" , 'Test_Accuracy', 'F1-Score', 'AUC_Score']

x.add_row([ "LR with BOW" , "l1" , 0.1 , 90.708 , 94.59 ,79.34 ])
x.add_row([ "LR with BOW without Standardization" , "l1" , 0.1 , 88.0 , 93.00 , 73.07 ])

x.add_row([ "LR with BOW" , "l2" , 0.001 , 91.268 ,94.974, 78.00 ])
x.add_row([ "LR with TFIDF" , "l1" ,0.1 , 91.444 , 95.013 ,81.821 ])
x.add_row([ "LR with TFIDF" , "l2" , 0.0001 , 91.776 , 95.322,76.296 ])

x.add_row([ "LR with AVG_W2V" , "l1" , 0.1, 88.74 , 93.62 ,69.338 ])
x.add_row([ "LR with AVG_W2V" , "l2" , 0.1 , 88.764 , 93.632,69.477])
```

```
x.add_row([ "LR with AVG_TFIDF" , "l1" , 0.5 , 87.527 , 93.000, 64.314
])
x.add_row([ "LR with AVG_TFIDF" , "l2" ,0.1, 87.539 , 93.013 , 64.33
])
```

```
print(x)
```

```
+-----+-----+-----+
+-----+-----+-----+
| LogisticRegression with Different Vectorization | penalty | C-value |
| Test_Accuracy | F1-Score | AUC_Score |
+-----+-----+-----+
+-----+-----+-----+
|          LR with BOW          | l1 | 0.1 |
| 90.708 | 94.59 | 79.34 |
| LR with BOW without Standardization | l1 | 0.1 |
| 88.0 | 93.0 | 73.07 |
|          LR with BOW          | l2 | 0.001 |
| 91.268 | 94.974 | 78.0 |
|          LR with TFIDF          | l1 | 0.1 |
| 91.444 | 95.013 | 81.821 |
|          LR with TFIDF          | l2 | 0.0001 |
| 91.776 | 95.322 | 76.296 |
|          LR with AVG_W2V          | l1 | 0.1 |
| 88.74 | 93.62 | 69.338 |
|          LR with AVG_W2V          | l2 | 0.1 |
| 88.764 | 93.632 | 69.477 |
|          LR with AVG_TFIDF          | l1 | 0.5 |
| 87.527 | 93.0 | 64.314 |
|          LR with AVG_TFIDF          | l2 | 0.1 |
| 87.539 | 93.013 | 64.33 |
+-----+-----+-----+
+-----+-----+-----+
```

Summy:

- Have Applied LR with All vectors and with both penalty ie L1 and L2
- Among all vectorization TFIDF with penalty L1 gives high accuracy
- Done LR with Standardization Data and without Standardization Data what i have observed is there is good amount of difference in both of the outputs
- StandardizationData output performs much better than not StandardizationData in LR
- Have Done Perbutation test on LR with BOW and observed 10's percentile on weight_difference
- TFIDF top 10Features :
 - top 10 negative features = [disappoint , veri disappoint , worst , terribl , return , horribl , th rew , wont buy , wast money]
 - top 10 Postive features = [great , love , best , good , delici , perfect , excel , favorit , nice , wonder]

In [0]: